

A Generalized Focused Information Criterion for GMM*

Minsu Chang Francis J. DiTraglia[†]

University of Pennsylvania

This Version: September 7, 2017 First Version: February 15, 2013

Abstract

This paper proposes a criterion for simultaneous GMM model and moment selection: the generalized focused information criterion (GFIC). Rather than attempting to identify the “true” specification, the GFIC chooses from a set of potentially misspecified moment conditions and parameter restrictions to minimize the mean-squared error (MSE) of a user-specified target parameter. The intent of the GFIC is to formalize a situation common in applied practice. An applied researcher begins with a set of fairly weak “baseline” assumptions, assumed to be correct, and must decide whether to impose any of a number of stronger, more controversial “suspect” assumptions that yield parameter restrictions, additional moment conditions, or both. Provided that the baseline assumptions identify the model, we show how to construct an asymptotically unbiased estimator of the asymptotic MSE to select over these suspect assumptions: the GFIC. We go on to provide results for post-selection inference and model averaging that can be applied both to the GFIC and various alternative selection criteria. We specialize the GFIC to three panel data examples: choosing between random and fixed effects estimators, selecting over exogeneity assumptions and lag lengths in a dynamic panel model, and choosing between pooled and mean-group estimators of an average effect in the presence of heterogeneity. In the random versus fixed effects example, we also propose an averaging estimator that optimally weights the two estimators. The GFIC performs well in simulations. We conclude by applying the GFIC to a dynamic panel data model for the price elasticity of cigarette demand.

Keywords: Model Selection, Moment selection, Model averaging, Panel Data, GMM Estimation, Focused Information Criterion, Post-selection estimators

JEL Codes: C23, C52

*We thank Manuel Arellano, Otilia Boldea, Camilo García-Jimeno, Bruce Hansen, Frank Kleibergen, and seminar participants at the 2013 Latin American Workshop in Econometrics, the 2014 Midwest Econometrics Group meetings, Tilburg, the Tinbergen Institute, and the University of Wisconsin for helpful comments.

[†]Corresponding Author: fditra@sas.upenn.edu, 3718 Locust Walk, Philadelphia, PA 19104

1 Introduction

An econometric model is a tool for answering a particular research question: different questions may suggest different models for the same data. And the fact that a model is wrong, as the old saying goes, does not prevent it from being useful. This paper proposes a novel selection criterion for GMM estimation that takes both of these points to heart: the generalized focused information criterion (GFIC). Rather than attempting to identify the correct specification, the GFIC chooses from a set of potentially mis-specified moment conditions and parameter restrictions to yield the smallest mean squared error (MSE) estimator of a user-specified scalar target parameter. We derive the GFIC under local mis-specification, using asymptotic mean squared error (AMSE) to approximate finite-sample MSE. In this framework mis-specification, while present for any fixed sample size, disappears in the limit so that asymptotic variance and squared bias remain comparable. GMM estimators remain consistent under local mis-specification but their limit distributions show an asymptotic bias. Adding an additional moment condition or imposing a parameter restriction generally reduces asymptotic variance but, if incorrectly specified, introduces a source of bias. The GFIC trades off these two effects in the first-order asymptotic expansion of an estimator to approximate its finite sample behavior.

The GFIC takes its motivation from a situation that is common in empirical practice. A researcher who hopes to estimate a parameter of interest μ must decide which assumptions to use. On the one hand is a set of relatively uncontroversial “baseline” assumptions. We suppose that the baseline assumptions are correct and identify μ . But the very fact that the baseline assumptions do not raise eyebrows suggests that they may not be especially informative about μ . On the other hand are one or more stronger controversial “suspect” assumptions. These stronger assumptions are expected to be much more informative about μ . If we were certain that they were correct, we would definitely choose to impose them in estimation. Indeed, by continuity, even if they were *nearly* correct, imposing the suspect assumptions could yield a favorable bias-variance tradeoff. This is the essential idea behind the GFIC. When the baseline assumptions identify the model, the GFIC provides an asymptotically unbiased estimator of AMSE.

The GFIC is an extension of the focused moment selection criterion (FMSC) of DiTraglia (2016). While the FMSC considers the problem of selecting moment conditions while holding the model specification *fixed*, the GFIC allows us to select over both aspects of our specification simultaneously. This extension is particularly valuable in panel data applications, where we may, for example, wish to carry out selection over the lag specification as well as the exogeneity assumptions used to estimate a dynamic panel model. We specialize the GFIC

to a dynamic panel example below, in addition to two others: selecting between random and fixed effects estimators, and choosing between pooled OLS and mean-group estimators of an average effect in the presence of heterogeneity. In addition to extending the FMSC to a broader class of problems and deriving specific results for well-known panel data problems, we also extend the results of DiTraglia (2016) on post-selection and moment-average estimators to the more general setting of the GFIC. In the random versus fixed effects example, we further derive an averaging estimator that optimally combines the information contained in the two specifications.¹ The GFIC and averaging estimator perform well in simulation studies for our examples. We conclude with an empirical example applying the GFIC to a panel data model for the price elasticity of cigarette demand.

As its name suggests, the GFIC is related to the focused information criterion (FIC) of Claeskens and Hjort (2003), a model selection procedure for maximum likelihood estimators that uses local mis-specification to approximate the MSE of a target parameter. Like the FIC and related proposals, e.g. Schorfheide (2005), the GFIC uses local mis-specification to derive a risk-based selection criterion. Unlike them, however, the GFIC provides both moment and model selection for general GMM estimators. If the moment conditions used in estimation are the score of a maximum likelihood model and we consider model selection only, then the GFIC reduces to the FIC. Thus, the GFIC extends both the FIC and the FMSC of DiTraglia (2016). Comparatively few papers propose criteria for simultaneous GMM model and moment selection under mis-specification.² Andrews and Lu (2001) propose a family of selection criteria by adding appropriate penalty and “bonus” terms to the J-test statistic, yielding analogues of AIC, BIC, and the Hannan-Quinn information criterion. Hong et al. (2003) extend this idea to generalized empirical likelihood (GEL). The principal goal of both papers is consistent selection: they state conditions under which the correct model and all correct moment conditions are chosen in the limit. As a refinement to this approach, Lai et al. (2008) suggest a two-step procedure: first consistently eliminate incorrect models using an empirical log-likelihood ratio criterion, and then select from the remaining models using a bootstrap covariance matrix estimator. The point of the second step is to address a shortcoming in the standard limit theory. While first-order asymptotic efficiency requires that we use all available correctly specified moment conditions, this can lead to a deterioration in finite sample performance if some conditions are only weakly informative. Hall and Peixe (2003) make a similar point about the dangers of including “redundant” moment conditions while Caner (2009) proposes a lasso-type GMM estimator to consistently

¹Although we are unaware of any other proposal for averaging fixed and random effects estimator, our general approach to averaging is shared by Hjort and Claeskens (2003) and is related to Hansen (2017).

²See Smith (1992) for an approach to GMM model selection based on non-nested hypothesis testing. For a detailed discussion of the literature on moment selection, see DiTraglia (2016).

remove redundant parameters.

In contrast to these suggestions, the GFIC does not aim to identify the correct model and moment conditions: its goal is a low MSE estimate of a quantity of interest, even if this entails using a specification that is not exactly correct. As such, the GFIC is an “efficient” rather than a consistent selection criterion. There is an unavoidable trade-off between consistent selection and estimators with desirable risk properties (Yang, 2005). Indeed, the worst-case risk of any consistent selection procedure is *unbounded* as sample size tends to infinity (Leeb and Pötscher, 2008). In this sense, the fact that the GFIC is not consistent is a benefit rather than a liability. As we show in simulations below, its worst-case performance is much better than that of competing selection procedures.

Although not strictly a selection procedure, the combined moments (CM) estimator of Judge and Mittelhammer (2007) takes a similar perspective to that of the GFIC, emphasizing that incorporating the information from an incorrect specification could lead to a favorable bias-variance tradeoff under the right circumstances. Their proposal uses a Cressie-Read divergence measure to combine the information from competing moment specifications, for example OLS versus two-stage least squares, yielding a data-driven compromise estimator. Unlike the GFIC, however, the CM estimator is not targeted to a particular research goal. A key point of the GFIC is that two different researchers working with the same dataset may find it optimal to use different specifications. We explore this idea further in our dynamic panel example below.

The remainder of this paper is organized as follows. Section 2 derives the asymptotic distribution of GMM estimators under locally mis-specified moment conditions and parameter restrictions. Section 3 uses this information to calculate the AMSE of a user-specified target parameter and provides asymptotically unbiased estimators of the required bias parameters, yielding the GFIC. Section 4 extends the results on averaging estimators and post-selection inference from DiTraglia (2016) to the more general setting of this paper. Section 5 specializes the GFIC to three panel data examples, and Section 6 presents simulation results. Finally, Section 7 presents our empirical example and Section 8 concludes. Proofs and supplementary simulation results appear in the Appendix.

2 Asymptotic Framework

Let $f(\cdot, \cdot)$ be a $(p + q)$ -vector of moment functions of a random vector Z and an $(r + s)$ -dimensional parameter vector β . To represent moment selection, we partition the moment functions according to $f(\cdot, \cdot) = (g(\cdot, \cdot)', h(\cdot, \cdot)')'$ where $g(\cdot, \cdot)$ and $h(\cdot, \cdot)$ are p - and q -vectors. The moment condition associated with $g(\cdot, \cdot)$ is assumed to be correct, while that associated

with $h(\cdot, \cdot)$ is locally mis-specified. The moment selection problem is to choose which, if any, of the elements of h to use in estimation. To represent model selection, we partition the full parameter vector according to $\beta = (\gamma', \theta')'$, where γ is an r -vector and θ an s -vector of parameters. The model selection problem is to decide which if any of the elements of γ to estimate, and which to set equal to the corresponding elements of γ_0 , an r -vector of known constants. The parameters contained in θ are those that we always estimate, the “protected” parameters. Any specification that does not estimate the full parameter vector β is locally mis-specified. The precise form of the local mis-specification, over parameter restrictions and moment conditions, is as follows.

Assumption 2.1 (Local Mis-specification). *Let $\{Z_{ni}: 1 \leq i \leq n, n = 1, 2, \dots\}$ be an iid triangular array of random vectors defined on a probability space $(\Upsilon, \mathcal{F}, P)$ satisfying*

- (a) $E[g(Z_{ni}, \gamma_n, \theta_0)] = 0$, with $\gamma_n = \gamma_0 + n^{-1/2}\delta$
- (b) $E[h(Z_{ni}, \gamma_n, \theta_0)] = \tau_n$, with $\tau_n = n^{-1/2}\tau$
- (c) $\{f(Z_{ni}, \gamma_n, \theta_0): 1 \leq i \leq n, n = 1, 2, \dots\}$ is uniformly integrable, and
- (d) $Z_{ni} \xrightarrow{d} Z_i$.

where γ_0 is a (known) r -vector of parameter restrictions, δ an unknown r -vector of constants, and τ an unknown q -vector of constants.

For simplicity, and because it is the case for all examples we consider below, we assume that the triangular array from Assumption 2.1 is iid, although this is not strictly necessary.

Under Assumption 2.1, the true parameter vector $\beta_n = (\gamma'_n, \theta'_0)'$, changes with sample size but converges to $\beta_0 = (\gamma'_0, \theta'_0)'$ as $n \rightarrow \infty$. Unless some elements of δ are zero, any estimator that restricts γ is mis-specified for fixed n . In the limit, however, the restriction $\gamma = \gamma_0$ holds. Similarly, for any fixed sample size n , the expectation of h evaluated at the true parameter value β_n depends on the unknown constant vector τ , but this source of mis-specification disappears in the limit. Thus, under Assumption 2.1, only estimators that use moment conditions from g to estimate the full parameter vector β are correctly specified. In the limit, however, *every* estimator is correctly specified, regardless of which elements of γ it restricts and which elements of h it includes. The purpose of local mis-specification is to ensure that squared asymptotic bias is of the same order as asymptotic variance: Assumption 2.1 is a device rather than literal description of real-world data. Note that, by Assumption 2.1, the limiting random variable Z_i satisfies the population moment

condition $E[f(Z_i, \gamma_0, \theta_0)] = 0$. Since the Z_i are assumed to have a common marginal law, we will use the shorthand Z for Z_i throughout. Accordingly, define:

$$F = \begin{bmatrix} \nabla_{\gamma'} g(Z, \gamma_0, \theta_0) & \nabla_{\theta'} g(Z, \gamma_0, \theta_0) \\ \nabla_{\gamma'} h(Z, \gamma_0, \theta_0) & \nabla_{\theta'} h(Z, \gamma_0, \theta_0) \end{bmatrix} \quad (1)$$

partitioned according to

$$F = \begin{bmatrix} F_\gamma & F_\theta \end{bmatrix} = \begin{bmatrix} G_\gamma & G_\theta \\ H_\gamma & H_\theta \end{bmatrix} = \begin{bmatrix} G \\ H \end{bmatrix} \quad (2)$$

and define

$$\Omega = Var \begin{bmatrix} g(Z, \gamma_0, \theta_0) \\ h(Z, \gamma_0, \theta_0) \end{bmatrix} = \begin{bmatrix} \Omega_{gg} & \Omega_{gh} \\ \Omega_{hg} & \Omega_{hh} \end{bmatrix}. \quad (3)$$

Each of these expressions involves the limiting random variable Z rather than Z_{ni} . Thus, the corresponding expectations are taken with respect to a distribution for which all moment conditions have expectation zero evaluated at (γ_0, θ_0) . Finally, let $F(b, c) = \Xi_c F \Xi_b'$ and $\Omega_c = \Xi_c \Omega \Xi_c'$ and $W_c = \Xi_c W \Xi_c'$ where W is the positive definite probability limit of \widetilde{W} .

Before defining the estimators under consideration, we require some further notation. Let b be a *model selection vector*, an r -vector of ones and zeros indicating which elements of γ we have chosen to estimate. When $b = 1_r$, where 1_m represents an m -vector of ones, we estimate both θ and the full vector γ . When $b = 0_r$, where 0_m denotes an m -vector of zeros, we estimate only θ , setting $\gamma = \gamma_0$. More generally, we estimate $|b|$ components of γ and set the others equal to the corresponding elements of γ_0 . Let $\gamma^{(b)}$ be the $|b|$ -dimensional subvector of γ corresponding to those elements selected for estimation. Similarly, let $\gamma_0^{(-b)}$ denote the $(r - |b|)$ -dimensional subvector containing the values to which we set those components of γ that are *not* estimated. Analogously, let $c = (c'_g, c'_h)'$ be a *moment selection vector*, a $(p + q)$ -vector of ones and zeros indicating which of the moment conditions we have chosen to use in estimation. We denote by $|c|$ the total number of moment conditions used in estimation. Let \mathcal{BC} denote the collection of all model and moment selection pairs (b, c) under consideration.

To express moment and model selection in matrix form, we define the selection matrices Ξ_b and Ξ_c . Multiplying β by the $(|b| + s) \times (r + s)$ *model selection matrix* Ξ_b extracts the elements corresponding to θ and the subset of γ indicated by the model selection vector b . Thus $\Xi_b \beta = (\gamma^{(b)'}, \theta')'$. Similarly, multiplying a vector by the $|c| \times (p + q)$ moment selection matrix Ξ_c extracts the components corresponding to the moment conditions indicated by the moment selection vector c .

To express the estimators themselves, define the sample analogue of the expectations in

Assumption 2.1 as follows,

$$f_n(\beta) = \frac{1}{n} \sum_{i=1}^n f(Z_{ni}, \gamma, \theta) = \begin{bmatrix} g_n(\beta) \\ h_n(\beta) \end{bmatrix} = \begin{bmatrix} n^{-1} \sum_{i=1}^n g(Z_{ni}, \theta, \gamma) \\ n^{-1} \sum_{i=1}^n h(Z_{ni}, \theta, \gamma) \end{bmatrix} \quad (4)$$

and let \widetilde{W} be a $(p+q) \times (p+q)$ positive semi-definite weighting matrix

$$\widetilde{W} = \begin{bmatrix} \widetilde{W}_{gg} & \widetilde{W}_{gh} \\ \widetilde{W}_{hg} & \widetilde{W}_{hh} \end{bmatrix} \quad (5)$$

partitioned conformably to the partition of $f(Z, \beta)$ by $g(Z, \beta)$ and $h(Z, \beta)$. Each model and moment selection pair $(b, c) \in \mathcal{BC}$ defines a $(|b| + s)$ -dimensional estimator $\widehat{\beta}(b, c) = (\widehat{\gamma}^{(b)}(b, c)', \widehat{\theta}(b, c)')'$ of $\beta^{(b)} = (\gamma^{(b)'}, \theta')'$ according to

$$\widehat{\beta}(b, c) = \arg \min_{\beta^{(b)} \in \mathbf{B}^{(b)}} \left[\Xi_c f_n \left(\beta^{(b)}, \gamma_0^{(-b)} \right) \right]' \left[\Xi_c \widetilde{W} \Xi_c' \right] \left[\Xi_c f_n \left(\beta^{(b)}, \gamma_0^{(-b)} \right) \right]. \quad (6)$$

We now state a number of standard high-level regularity conditions that will be assumed throughout our derivations below.

Assumption 2.2 (High-level Regularity Conditions).

- (a) β_0 lies in the interior of Θ , a compact set
- (b) $\widetilde{W} \rightarrow_p W$, a positive definite matrix
- (c) $W_c \Xi_c E[f(Z, \beta)] = 0$ if and only if $\beta = \beta_0$, where $W_c = \Xi_c W \Xi_c'$
- (d) $E[f(Z, \beta)]$ is continuous on Θ
- (e) $\sup_{\beta \in \Theta} \|f_n(\beta) - E[f(Z, \beta)]\| \rightarrow_p 0$
- (f) f is Z -almost surely differentiable in an open neighborhood \mathcal{B} of β_0
- (g) $\sup_{\beta \in \Theta} \|\nabla_\beta f_n(\beta) - F(\beta)\| \rightarrow_p 0$
- (h) $\sqrt{n}f_n(\beta_0) - \sqrt{n}E[f(Z, \beta_0)] \rightarrow_d \mathcal{N}$ where $\mathcal{N} \sim N_{p+q}(0, \Omega)$
- (i) $F(b, c)'W_c F(b, c)$ is invertible, where $W_c = \Xi_c W \Xi_c'$

A particularly important special case is the estimator using only the moment conditions in g to estimate the full parameter vector $\beta = (\gamma', \theta')'$. We call this the *valid* estimator and denote it by $\widehat{\beta}_v$. Because it is assumed to be correctly specified both for finite n and in

the limit, the valid estimator contains the information we use to identify τ and δ , and thus carry out moment and model selection. We assume throughout that the valid estimator is identified.

Assumption 2.3 (GFIC Identification Condition). *Let $\widehat{\beta}_v$ denote the GMM estimator for $\beta = (\gamma', \theta')'$ based solely on the moment conditions contained in g :*

$$\widehat{\beta}_v = \begin{bmatrix} \widehat{\gamma}_v \\ \widehat{\theta}_v \end{bmatrix} = \arg \min_{\beta \in \mathbf{B}} g_n(\beta)' \widetilde{W}_{gg} g_n(\beta). \quad (7)$$

We call this the “valid estimator” and assume that it satisfies all the conditions of Assumption 2.2. In particular assume that $\widehat{\beta}_v$ is identified, which implies $p \geq r + s$.

Because Assumption 2.1 ensures that they are correctly specified in the limit, *all* candidate specifications $(b, c) \in \mathcal{BC}$ provide consistent estimators of θ_0 under standard, high level regularity conditions (see Assumption 2.2). Essential differences arise, however, when we consider their respective asymptotic distributions. Under Assumption 2.1, both δ and τ induce a bias term in the limiting distribution of $\sqrt{n} \left(\widehat{\beta}(b, c) - \beta_0^{(b)} \right)$. The key result is as follows.

Theorem 2.1 (Asymptotic Distribution). *Under Assumptions 2.1–2.2*

$$\sqrt{n} \left(\widehat{\beta}(b, c) - \beta_0^{(b)} \right) \xrightarrow{d} -K(b, c) \Xi_c \left(\mathcal{N} + \begin{bmatrix} 0 \\ \tau \end{bmatrix} - F_\gamma \delta \right) \quad (8)$$

where $\beta_0^{(b)'} = (\gamma_0^{(b)'}, \theta_0')'$, $K(b, c) = [F(b, c)' W_c F(b, c)]^{-1} F(b, c)' W_c$ and

$$\mathcal{N} = \begin{bmatrix} \mathcal{N}_g \\ \mathcal{N}_h \end{bmatrix} \sim N \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} \Omega_{gg} & \Omega_{gh} \\ \Omega_{hg} & \Omega_{hh} \end{bmatrix} \right). \quad (9)$$

Because it employs the correct specification, the valid estimator of θ shows no asymptotic bias. Moreover, the valid estimator of γ has an asymptotic distribution that is centered around δ , suggesting an estimator of this bias parameter.

Corollary 2.1 (Asymptotic Distribution of Valid Estimator). *Under Assumptions 2.1–2.2*

$$\sqrt{n} \left(\widehat{\beta}_v - \beta_0 \right) = \sqrt{n} \begin{pmatrix} \widehat{\theta}_v - \theta_0 \\ \widehat{\gamma}_v - \gamma_0 \end{pmatrix} \xrightarrow{d} \begin{bmatrix} 0 \\ \delta \end{bmatrix} - K_v \mathcal{N}_g$$

where $K_v = [G' W_{gg} G]^{-1} G' W_{gg}$ and $W_{gg} = \text{plim } \widetilde{W}_{gg}$.

We use these results in the following section to construct the GFIC.

3 The GFIC

The GFIC chooses among potentially incorrect moment conditions and parameter restrictions to minimize estimator AMSE for a scalar target parameter. Denote this target parameter by $\mu = \varphi(\gamma, \theta)$, where φ is a real-valued, almost surely continuous function of the underlying model parameters θ and γ . Let $\mu_n = \varphi(\gamma_n, \theta_0)$ and define μ_0 and $\hat{\mu}(b, c)$ analogously. By Theorem 2.1 and the delta method, we have the following result.

Corollary 3.1. *Under the hypotheses of Theorem 2.1,*

$$\sqrt{n}(\hat{\mu}(b, c) - \mu_0) \xrightarrow{d} -\nabla_{\beta}\varphi'_0\Xi'_bK(b, c)\Xi_c \left(\mathcal{N} + \begin{bmatrix} 0 \\ \tau \end{bmatrix} - F_{\gamma}\delta \right)$$

where $\varphi_0 = \varphi(\gamma_0, \theta_0)$.

The true value of μ , however, is μ_n rather than μ_0 under Assumption 2.1. Accordingly, to calculate AMSE we recenter the limit distribution as follows.

Corollary 3.2. *Under the hypotheses of Theorem 2.1,*

$$\sqrt{n}(\hat{\mu}(b, c) - \mu_n) \xrightarrow{d} -\nabla_{\beta}\varphi'_0\Xi'_bK(b, c)\Xi_c \left(\mathcal{N} + \begin{bmatrix} 0 \\ \tau \end{bmatrix} - F_{\gamma}\delta \right) - \nabla_{\gamma}\varphi'_0\delta$$

where $\varphi_0 = \varphi(\gamma_0, \theta_0)$.

We see that the limiting distribution of $\hat{\mu}(b, c)$ is not, in general, centered around zero: both τ and δ induce an asymptotic bias. Note that, while τ enters the limit distribution only once, δ has two distinct effects. First, like τ , it shifts the limit distribution of $\sqrt{n}f_n(\gamma_0, \theta_0)$ away from zero, thereby influencing the asymptotic behavior of $\sqrt{n}(\hat{\mu}(b, c) - \mu_0)$. Second, unless the derivative of φ with respect to γ is zero at (γ_0, θ_0) , δ induces a second source of bias when $\hat{\mu}(b, c)$ is recentered around μ_n . Crucially, this second source of bias exactly cancels the asymptotic bias present in the limit distribution of $\hat{\gamma}_v$. Thus, the valid estimator of μ is asymptotically unbiased and its AMSE equals its asymptotic variance.

Corollary 3.3. *Under the hypotheses of Theorem 2.1,*

$$\sqrt{n}(\hat{\mu}_v - \mu_n) \xrightarrow{d} -\nabla_{\beta}\varphi(\gamma_0, \theta_0)'K_v\mathcal{N}_g$$

where $\hat{\mu}_v = \varphi(\hat{\gamma}_v, \hat{\theta}_v)$. Thus, the valid estimator $\hat{\mu}_v$ shows no asymptotic bias and has asymptotic variance $\nabla_{\beta}\varphi(\gamma_0, \theta_0)'K_v\Omega_{gg}K'_v\nabla_{\beta}\varphi(\gamma_0, \theta_0)$.

Using Corollary 3.2, the AMSE of $\hat{\mu}(b, c)$ is as follows,

$$\text{AMSE}(\hat{\mu}(b, c)) = \text{AVAR}(\hat{\mu}(b, c)) + \text{BIAS}(\hat{\mu}(b, c))^2 \quad (10)$$

$$\text{AVAR}(\hat{\mu}(b, c)) = \nabla_{\beta} \varphi_0' \Xi_b' K(b, c) \Omega_c K(b, c)' \Xi_b \nabla_{\beta} \varphi_0 \quad (11)$$

$$\text{BIAS}(\hat{\mu}(b, c)) = -\nabla_{\beta} \varphi_0' M(b, c) \begin{bmatrix} \delta \\ \tau \end{bmatrix} \quad (12)$$

where

$$M(b, c) = \Xi_b' K(b, c) \Xi_c \begin{bmatrix} -G_{\gamma} & 0 \\ -H_{\gamma} & I \end{bmatrix} + \begin{bmatrix} I_r & 0_{r \times q} \\ 0_{p \times r} & 0_{s \times q} \end{bmatrix} \quad (13)$$

The idea behind the GFIC is to construct an estimate $\widehat{\text{AMSE}}(\hat{\mu}(b, c))$ and choose the specification $(b^*, c^*) \in \mathcal{BC}$ that makes this quantity as small as possible. As a side-effect of the consistency of the estimators $\hat{\beta}(b, c)$, the usual sample analogues provide consistent estimators of $K(b, c)$ and $F'_{\gamma} = (G'_{\gamma}, H'_{\gamma})$ under Assumption 2.1, and $\varphi(\hat{\theta}_v, \gamma_0)$ is consistent for φ_0 . Consistent estimators of Ω are also readily available under local mis-specification although the best choice may depend on the situation.³ Since γ_0 is known, as are Ξ_b and Ξ_c , only δ and τ remain to be estimated. Unfortunately, neither of these quantities is consistently estimable under local mis-specification. Intuitively, the data become less and less informative about τ and δ as the sample size increases since each term is divided by \sqrt{n} . Multiplying through by \sqrt{n} counteracts this effect, but also stabilizes the variance of our estimators. Hence, the best we can do is to construct *asymptotically unbiased* estimators of τ and δ . Corollary 3.4 provides the required estimator for δ , namely $\hat{\delta} = \sqrt{n}(\hat{\gamma}_v - \gamma_0)$, while Lemma 3.1 provides an asymptotically unbiased estimator of τ by plugging $\hat{\beta}_v$ into the sample analogue of the h -block of moment conditions.

Corollary 3.4. *Under the hypotheses of Theorem 2.1, $\hat{\delta} = \sqrt{n}(\hat{\gamma}_v - \gamma_0) \xrightarrow{d} \delta - K_v^{\gamma} \mathcal{N}_g$ where $K_v = [G'W_{gg}G]^{-1}G'W_{gg} = (K_v^{\gamma'}, K_v^{\theta'})'$, so $\hat{\delta}$ is an asymptotically unbiased estimator of δ .*

Lemma 3.1. *Under the hypotheses of Theorem 2.1,*

$$\hat{\tau} = \sqrt{n}h_n(\hat{\beta}_v) \xrightarrow{d} \tau - HK_v\mathcal{N}_g + \mathcal{N}_h$$

where $K_v = [G'W_{gg}G]^{-1}G'W_{gg}$. Hence, $\hat{\tau}$ is an asymptotically unbiased estimator of τ .

Combining Corollary 3.4 and Lemma 3.1, gives the joint distribution of $\hat{\delta}$ and $\hat{\tau}$.

³We discuss this in more detail for our examples in Section 5 below.

Theorem 3.1. *Under the hypotheses of Theorem 2.1,*

$$\begin{bmatrix} \widehat{\delta} \\ \widehat{\tau} \end{bmatrix} = \sqrt{n} \begin{bmatrix} (\widehat{\gamma}_v - \gamma_0) \\ h_n(\widehat{\beta}_v) \end{bmatrix} \xrightarrow{d} \begin{bmatrix} \delta \\ \tau \end{bmatrix} + \Psi \mathcal{N}, \quad \Psi = \begin{bmatrix} -K_v^\gamma & \mathbf{0} \\ -HK_v & I \end{bmatrix}$$

where $K_v = [G'W_{gg}G]^{-1}G'W_{gg}$ is partitioned according to $K_v' = (K_v^{\gamma'}, K_v^{\theta'})$.

Now, we see immediately from Equation 12 that

$$\text{BIAS}(\widehat{\mu}(b, c))^2 = \nabla_\beta \varphi_0' M(b, c) \begin{bmatrix} \delta\delta' & \delta\tau' \\ \tau\delta' & \tau\tau' \end{bmatrix} M(b, c)' \nabla_\beta \varphi_0$$

Thus, the bias parameters τ and δ enter the AMSE expression in Equation 10 as outer products: $\tau\tau'$, $\delta\delta'$ and $\tau\delta'$. Although $\widehat{\tau}$ and $\widehat{\delta}$ are asymptotically unbiased estimators of τ and δ , it does *not* follow that $\widehat{\tau}\widehat{\tau}$, $\widehat{\delta}\widehat{\delta}$ and $\widehat{\tau}\widehat{\delta}'$ are asymptotically unbiased estimators of $\tau\tau'$, $\delta\delta'$, and $\tau\delta'$. The following result shows how to adjust these quantities to provide the required asymptotically unbiased estimates.

Corollary 3.5. *Suppose that $\widehat{\Psi}$ and $\widehat{\Omega}$ are consistent estimators of Ψ and Ω . Then, \widehat{B} is an asymptotically unbiased estimator of B , where*

$$\widehat{B} = \begin{bmatrix} \widehat{\delta}\widehat{\delta}' & \widehat{\delta}\widehat{\tau}' \\ \widehat{\tau}\widehat{\delta}' & \widehat{\tau}\widehat{\tau}' \end{bmatrix} - \widehat{\Psi}\widehat{\Omega}\widehat{\Psi}', \quad B = \begin{bmatrix} \delta\delta' & \delta\tau' \\ \tau\delta' & \tau\tau' \end{bmatrix}.$$

Combining Corollary 3.5 with consistent estimates of the remaining quantities yields the GFIC, an asymptotically unbiased estimator of the AMSE of our estimator of a target parameter μ under each specification $(b, c) \in \mathcal{BC}$

$$\text{GFIC}(b, c) = \nabla_\beta \widehat{\varphi}_0' \left[\Xi_b' \widehat{K}(b, c) \widehat{\Omega}_c \widehat{K}(b, c)' \Xi_b + \widehat{M}(b, c) \widehat{B} \widehat{M}(b, c)' \right] \nabla_\beta \widehat{\varphi}_0. \quad (14)$$

We choose the specification (b^*, c^*) that minimizes the GFIC over the candidate set \mathcal{BC} .

4 Averaging and Post-Selection Inference

While we are primarily concerned in this paper with the mean-squared error performance of our proposed selection techniques, it is important to have tools for carrying out inference post-selection. In this section we briefly present results that can be used to carry out valid inference for a range of model averaging and post-selection estimators, including the GFIC.⁴

⁴We direct the reader to DiTraglia (2016) and the references contained therein for a more detailed discussion of inference post-selection.

The GFIC is an efficient rather than consistent selection criterion: it aims to estimate a particular target parameter with minimum AMSE rather than selecting the correct specification with probability approaching one in the limit. As pointed out by Yang (2005), among others, there is an unavoidable trade-off between consistent selection and desirable risk properties. Faced with this dilemma, the GFIC sacrifices consistency in the interest of low AMSE. Because it is not a consistent criterion, the GFIC remains random *even in the limit*. We can see this from Equation 14 in Section 3 and Corollary 3.5. While the quantities $\nabla_{\beta}\hat{\varphi}_0$, $\hat{K}(b, c)$, and $\hat{\Omega}_c$ are consistent estimators of their population counterparts, \hat{B} is only an asymptotically unbiased estimator of B and thus has a limiting distribution. In particular $\hat{B} \rightarrow_d \mathcal{B}(\mathcal{N}, \delta, \tau)$ where

$$\mathcal{B}(\mathcal{N}, \delta, \tau) = \left(\begin{bmatrix} \delta \\ \tau \end{bmatrix} + \Psi \mathcal{N} \right) \left(\begin{bmatrix} \delta \\ \tau \end{bmatrix} + \Psi \mathcal{N} \right)' - \Psi \Omega \Psi \quad (15)$$

Accordingly, to carry out inference post-GFIC, we need a limiting theory that is rich enough to accommodate *randomly-weighted* averages of the candidate estimators $\hat{\mu}(b, c)$. To this end, consider an estimator of the form

$$\hat{\mu} = \sum_{(b, c) \in \mathcal{BC}} \hat{\omega}(b, c) \hat{\mu}(b, c)$$

where $\hat{\mu}(b, c)$ denotes the target parameter under the moment conditions and parameter restrictions indexed by (b, c) , \mathcal{BC} denotes the full set of candidate specifications, and $\hat{\omega}(b, c)$ denotes a collection of data-dependent weights. These could be zero-one weights corresponding to a moment or model selection criterion, e.g. select the estimator that minimizes GFIC, or model averaging weights.⁵ We impose the following mild restrictions on the weights $\hat{\omega}$.

Assumption 4.1 (Conditions on the Weights).

- (a) $\sum_{(b, c) \in \mathcal{BC}} \hat{\omega}(b, c) = 1$, *almost surely*
- (b) For each $(b, c) \in \mathcal{BC}$, $\hat{\omega}(b, c) \xrightarrow{d} \psi(\mathcal{N}, \delta, \tau | b, c)$, a function of \mathcal{N} , δ , τ , and consistently estimable constants with at most countably many discontinuities.

Under the preceding conditions, we can derive the limit distribution of $\hat{\mu}$ shown in the following Corollary.

⁵For an example that averages over fixed and random effects estimators, see Section 5.1.

Corollary 4.1 (Limit Distribution of Averaging Estimators). *Under Assumption 4.1 and the hypotheses of Theorem 2.1, $\sqrt{n}(\hat{\mu} - \mu_n) \xrightarrow{d} \Lambda(\delta, \tau)$ where*

$$\Lambda(\delta, \tau) = -\nabla_{\beta} \varphi'_0 \sum_{(b,c) \in \mathcal{BC}} \psi(\mathcal{N}, \delta, \tau | b, c) \left\{ \Xi'_b K(b, c) \Xi_c \mathcal{N} + M(b, c) \begin{bmatrix} \delta \\ \tau \end{bmatrix} \right\} \quad (16)$$

Note that the limit distribution from the preceding corollary is highly non-normal: it is a *randomly* weighted average of a normal random vector, \mathcal{N} . To tabulate this distribution for the purposes of inference, we will in general need to resort to simulation. If τ and δ were known, the story would end here. In this case we could simply substitute consistent estimators of K and M and then repeatedly draw $\mathcal{N} \sim N(0, \hat{\Omega})$, where $\hat{\Omega}$ is a consistent estimator of Ω , to tabulate the distribution of Λ to arbitrary precision as follows.

Algorithm 4.1 (Approximating Quantiles of $\Lambda(\delta, \tau)$).

1. Generate J independent draws $\mathcal{N}_j \sim N(0, \hat{\Omega})$
2. Set $\Lambda_j(\delta, \tau) = -\nabla_{\beta} \hat{\varphi}'_0 \sum_{(b,c) \in \mathcal{BC}} \hat{\psi}(\mathcal{N}_j, \delta, \tau | b, c) \left\{ \Xi'_b \hat{K}(b, c) \Xi_c \mathcal{N}_j + \hat{M}(b, c) \begin{bmatrix} \delta \\ \tau \end{bmatrix} \right\}$
3. Using the $\Lambda_j(\delta, \tau)$, find $\hat{a}(\delta, \tau)$, $\hat{b}(\delta, \tau)$ so that $P \left\{ \hat{a}(\delta, \tau) \leq \Lambda(\delta, \tau) \leq \hat{b}(\delta, \tau) \right\} = 1 - \alpha$.

Unfortunately, no consistent estimators of τ or δ exist: all we have at our disposal are asymptotically unbiased estimators. The following “1-step” confidence interval is constructed by substituting these into Algorithm 4.1.

Algorithm 4.2 (1-Step Confidence Interval). Carry out of Algorithm 4.1 with τ and δ set equal to the estimators $\hat{\tau}$ and $\hat{\delta}$ from Theorem 3.1 to calculate $\hat{a}(\hat{\delta}, \hat{\tau})$ and $\hat{b}(\hat{\delta}, \hat{\tau})$. Then set $\text{CI}_1(\alpha) = \left[\hat{\mu} - \hat{b}(\hat{\delta}, \hat{\tau})/\sqrt{n}, \quad \hat{\mu} - \hat{a}(\hat{\delta}, \hat{\tau})/\sqrt{n} \right]$.

The 1-Step interval defined in Algorithm 4.2 is conceptually simple, easy to compute, and can perform well in practice.⁶ But as it fails to account for sampling uncertainty in $\hat{\tau}$, CI_1 does *not* necessarily yield asymptotically valid inference for μ . Fully valid inference requires the addition of a second step to the algorithm and comes at a cost: conservative rather than exact inference. In particular, the two-step procedure described in the following algorithm is guaranteed to yield an interval with asymptotic coverage probability of *no less than* $(1 - \alpha_1 - \alpha_2) \times 100\%$.

⁶For more discussion on this point, see DiTraglia (2016).

Algorithm 4.3 (2-Step Confidence Interval for $\hat{\mu}$).

1. Construct \mathcal{R} , a $(1 - \alpha_1) \times 100\%$ joint confidence region for (δ, τ) using Theorem 3.1.
2. For each $(\delta^*, \tau^*) \in \mathcal{R}$ carry out Algorithm 4.1 with $\alpha = \alpha_2$ yielding a $(1 - \alpha_2) \times 100\%$ confidence interval $[\hat{a}(\delta^*, \tau^*), \hat{b}(\delta^*, \tau^*)]$ for $\Lambda(\delta^*, \tau^*)$.
3. Set $\hat{a}_{min} = \min_{(\delta^*, \tau^*) \in \mathcal{R}} \hat{a}(\delta^*, \tau^*)$ and $\hat{b}_{max} = \max_{(\delta^*, \tau^*) \in \mathcal{R}} \hat{b}(\delta^*, \tau^*)$.
4. Construct the interval $CI_2(\alpha_1, \alpha_2) = [\hat{\mu} - \hat{b}_{max}/\sqrt{n}, \hat{\mu} - \hat{a}_{min}/\sqrt{n}]$.

Theorem 4.1 (2-Step Confidence Interval for $\hat{\mu}$). *Let $\nabla_{\beta}\hat{\varphi}_0$, $\hat{\psi}(\cdot|b, c)$, $\hat{K}(b, c)$ and $\hat{M}(b, c)$ be consistent estimators of $\nabla_{\beta}\varphi_0$, $\psi(\cdot|b, c)$, $K(b, c)$ and $M(b, c)$ and let R be a $(1 - \alpha_1) \times 100\%$ confidence region for (δ, τ) constructed from Theorem 3.1. Then $CI_2(\alpha_1, \alpha_2)$, defined in Algorithm 4.3 has asymptotic coverage probability no less than $1 - (\alpha_1 + \alpha_2)$ as $J, n \rightarrow \infty$.*

5 Examples of the GFIC in Panel Data Models

5.1 Random Effects versus Fixed Effects Example

In this section we consider a simple example in which the GFIC is used to choose between and average over alternative assumptions about individual heterogeneity: Random Effects versus Fixed Effects. For simplicity we consider the homoskedastic case and assume that any strictly exogenous regressors, including a constant term, have been “projected out” so we may treat all random variables as mean zero. To avoid triple subscripts in the notation, we further suppress the dependence of random variables on the cross-section dimension n except within statements of theorems. Suppose that

$$y_{it} = \beta x_{it} + v_{it} \tag{17}$$

$$v_{it} = \alpha_i + \varepsilon_{it} \tag{18}$$

for $i = 1, \dots, n$, $t = 1, \dots, T$ where ε_{it} is iid across i, t with $Var(\varepsilon_{it}) = \sigma_{\varepsilon}^2$ and α_i is iid across i with $Var(\alpha_i) = \sigma_{\alpha}^2$. Stacking observations for a given individual over time in the usual way, let $\mathbf{y}_i = (y_{i1}, \dots, y_{iT})'$ and define $\mathbf{x}_i, \mathbf{v}_i$ and $\boldsymbol{\varepsilon}_i$ analogously. Our goal in this example is to estimate β , the effect of x on y . Although x_{it} is uncorrelated with the time-varying portion of the error term, $Cov(x_{it}, \varepsilon_{it}) = 0$, we are unsure whether or not it is correlated

with the individual effect α_i . If we knew for certain that $Cov(x_{it}, \alpha_i) = 0$, we would prefer to report the “random effects” generalized least squares (GLS) estimator given by

$$\hat{\beta}_{GLS} = \left(\sum_{i=1}^n \mathbf{x}'_i \hat{\Omega}^{-1} \mathbf{x}_i \right)^{-1} \left(\sum_{i=1}^n \mathbf{x}'_i \hat{\Omega}^{-1} \mathbf{y}_i \right) \quad (19)$$

where $\hat{\Omega}^{-1}$ is a preliminary consistent estimator of

$$\Omega^{-1} = [Var(\mathbf{v}_i)]^{-1} = \frac{1}{\sigma_\epsilon^2} \left[I_T - \frac{\sigma_\alpha^2}{(T\sigma_\alpha^2 + \sigma_\epsilon^2)} \boldsymbol{\iota}_T \boldsymbol{\iota}'_T \right] \quad (20)$$

and I_T denotes the $T \times 1$ identity matrix and $\boldsymbol{\iota}_T$ a T -vector of ones. This estimator makes efficient use of the variation between and within individuals, resulting in an estimator with a lower variance. When $Cov(x_{it}, \alpha_i) \neq 0$, however, the random effects estimator is biased. Although its variance is higher than that of the GLS estimator, the “fixed effects” estimator given by

$$\hat{\beta}_{FE} = \left(\sum_{i=1}^n \mathbf{x}'_i Q \mathbf{x}_i \right)^{-1} \left(\sum_{i=1}^n \mathbf{x}'_i Q \mathbf{y}_i \right), \quad (21)$$

where $Q = I_T - \boldsymbol{\iota}_T \boldsymbol{\iota}'_T / T$, remains unbiased even when x_{it} is correlated with α_i .

The conventional wisdom holds that one should use the fixed effects estimator whenever $Cov(x_{it}, \alpha_i) \neq 0$. If the correlation between the regressor of interest and the individual effect is *sufficiently small*, however, the lower variance of the random effects estimator could more than compensate for its bias in a mean-squared error sense. This is precisely the possibility that we consider here using the GFIC. In this example, the local mis-specification assumption takes the form

$$\sum_{t=1}^T E[x_{it} \alpha_i] = \frac{\tau}{\sqrt{n}} \quad (22)$$

where τ is fixed, unknown constant. In the limit the random effects assumption that $Cov(x_{it}, \alpha_i) = 0$ holds, since $\tau/\sqrt{n} \rightarrow 0$. Unless $\tau = 0$, however, this assumption *fails* to hold for any finite sample size. An asymptotically unbiased estimator of τ for this example is given by

$$\hat{\tau} = (T\hat{\sigma}_\alpha^2 + \hat{\sigma}_\epsilon^2) \left[\frac{1}{\sqrt{n}} \sum_{i=1}^n \mathbf{x}'_i \hat{\Omega}^{-1} (\mathbf{y}_i - \mathbf{x}_i \hat{\beta}_{FE}) \right] \quad (23)$$

leading to the following result, from which we will construct the GFIC for this example.

Theorem 5.1 (Fixed versus Random Effects Limit Distributions). *Let $(\mathbf{x}_{ni}, \alpha_{ni}, \boldsymbol{\epsilon}_{ni})$ be an iid triangular array of random variables such that $Var(\boldsymbol{\epsilon}_i | \mathbf{x}_{ni}, \alpha_{ni}) \rightarrow \sigma_\epsilon^2 I_T$, $E[\mathbf{x}'_{ni} Q \boldsymbol{\epsilon}_{ni}] = 0$,*

and $E[\alpha_i \mathbf{l}'_T \mathbf{x}_{ni}] = \tau/\sqrt{n}$ for all n . Then, under standard regularity conditions,

$$\begin{bmatrix} \sqrt{n}(\hat{\beta}_{RE} - \beta) \\ \sqrt{n}(\hat{\beta}_{FE} - \beta) \\ \hat{\tau} \end{bmatrix} \xrightarrow{d} \left(\begin{bmatrix} c\tau \\ 0 \\ \tau \end{bmatrix}, \begin{bmatrix} \eta^2 & \eta^2 & 0 \\ \eta^2 & c^2\sigma^2 + \eta^2 & -c\sigma^2 \\ 0 & -c\sigma^2 & \sigma^2 \end{bmatrix} \right)$$

where $\eta^2 = E[\mathbf{x}'_i \Omega^{-1} \mathbf{x}_i]$, $c = E[\mathbf{x}'_i Q \mathbf{x}_i]/(T\sigma_\alpha^2 + \sigma_\varepsilon^2)$, and

$$\sigma^2 = \frac{(T\sigma_\alpha^2 + \sigma_\varepsilon^2)^2}{E[\mathbf{x}'_i \Omega^{-1} \mathbf{x}_i]} \left(\frac{\sigma_\varepsilon^2}{E[\mathbf{x}_i \Omega^{-1} \mathbf{x}_i] E[\mathbf{x}_i Q \mathbf{x}_i]} - 1 \right).$$

We see from Theorem 5.1 that $AMSE(\hat{\beta}_{RE}) = c^2\tau^2 + \eta^2$, $AMSE(\hat{\beta}_{FE}) = c^2\sigma^2 + \eta^2$, and $\hat{\tau}^2 - \sigma^2$ provides an asymptotically unbiased estimator of τ^2 . Thus, substituting $\hat{\tau}^2 - \sigma^2$ for τ and rearranging the preceding AMSE expressions, the GFIC tells us that we should select the random effects estimator whenever $|\hat{\tau}| \leq \sqrt{2}\sigma$. To implement this rule in practice, we construct a consistent estimator of σ^2 , for which we require estimators of $\sigma_\alpha^2, \sigma_\varepsilon^2$ and $\sigma_v^2 = Var(\alpha_i + \varepsilon_{it})$. We estimate these from the residuals

$$\hat{\epsilon}_{it} = (y_{it} - \bar{y}_i) - (x_{it} - \bar{x}_i)\hat{\beta}_{FE}; \quad \hat{v}_{it} = y_{it} - x_{it}\hat{\beta}_{OLS}$$

where $\hat{\beta}_{OLS}$ denotes the *pooled* OLS estimator of β , leading to the variance estimators

$$\hat{\sigma}_\alpha^2 = \hat{\sigma}_v^2 - \hat{\sigma}_\epsilon^2; \quad \hat{\sigma}_\epsilon^2 = \frac{1}{n(T-1)-1} \sum_{i=1}^n \sum_{t=1}^T \hat{\epsilon}_{it}^2; \quad \hat{\sigma}_v^2 = \frac{1}{nT-1} \sum_{i=1}^n \sum_{t=1}^T \hat{v}_{it}^2$$

Selection, of course, is a somewhat crude procedure: it is essentially an average that uses all-or-nothing weights. As a consequence, relatively small changes to the data could produce discontinuous changes in the weights, leading to a procedure with a high variance. Rather than selecting between the random effects and fixed effects estimators based on estimated AMSE, an alternative idea is to consider a more general weighted average of the form

$$\tilde{\beta}(\omega) = \omega\hat{\beta}_{FE} + (1-\omega)\hat{\beta}_{RE}$$

and for $\omega \in [0, 1]$ optimize the choice of ω to minimize AMSE. From Theorem 5.1 we see that the AMSE-minimizing value of ω is $\omega^* = (1 + \tau^2/\sigma^2)^{-1}$. Substituting our asymptotically unbiased estimator of τ^2 and our consistent estimator $\hat{\sigma}^2$ of σ^2 , we propose the following plug-in estimator of ω^*

$$\omega^* = \left[1 + \frac{\max\{\hat{\tau}^2 - \hat{\sigma}^2, 0\}}{\hat{\sigma}^2} \right]^{-1}$$

where we take the maximum over $\hat{\tau}^2 - \hat{\sigma}^2$ and zero so that $\hat{\omega}^*$ is between zero and one.

5.2 Dynamic Panel Example

We now specialize the GFIC to a dynamic panel model of the form

$$y_{it} = \theta x_{it} + \gamma_1 y_{it-1} + \cdots + \gamma_k y_{it-k} + \eta_i + v_{it} \quad (24)$$

where $i = 1, \dots, n$ indexes individuals and $t = 1, \dots, T$ indexes time periods. For simplicity, and without loss of generality, we suppose that there are no exogenous time-varying regressors and that all random variables are mean zero.⁷ The unobserved error η_i is a correlated individual effect: $\sigma_{x\eta} \equiv \mathbb{E}[x_{it}\eta_i]$ may not equal zero. The endogenous regressor x_{it} is assumed to be predetermined but not necessarily strictly exogenous: $\mathbb{E}[x_{it}v_{is}] = 0$ for all $s \geq t$ but may be nonzero for $s < t$. We assume throughout that y_{it} is stationary, which requires both x_{it} and u_{it} to be stationary and $|\boldsymbol{\gamma}| < 1$ where $\boldsymbol{\gamma} = (\gamma_1, \dots, \gamma_k)'$. Our goal is to estimate one of the following two target parameters with minimum MSE:

$$\mu_{SR} \equiv \theta, \quad \mu_{LR} \equiv \frac{\theta}{1 - (\gamma_1 + \cdots + \gamma_k)}. \quad (25)$$

where μ_{SR} denotes the short-run effect and μ_{LR} the long-run effect of x on y .

The question is which assumptions to use in estimation. Naturally, the answer may depend on whether our target is μ_{SR} or μ_{LR} . Our first decision is what assumption to impose on the relationship between x_{it} and v_{it} . This is the *moment selection* decision. We assumed above that x is predetermined. Imposing the stronger assumption of strict exogeneity gives us more and stronger moment conditions, but using these in estimation introduces a bias if x is not in fact strictly exogenous. Our second decision is how many lags of y to use in estimation. This is the *model selection* decision. The true model contains k lags of y . If we estimate only $r < k$ lags we not only have more degrees of freedom but more observations: every additional lag of y requires us to drop one time period from estimation. In the short panel datasets common in microeconomic applications, losing even one additional time period can represent a substantial loss of information. At the same time, unless $\gamma_{r+1} = \cdots = \gamma_k = 0$, failing to include all k lags in the model introduces a bias.

To eliminate the individual effects η_i we work in first differences. Defining Δ in the usual

⁷Alternatively, we can simply de-mean and project out any time-varying exogenous covariates after taking first-differences.

way, so that $\Delta y_{it} = y_{it} - y_{it-1}$ and so on, we can write Equation 24 as

$$\Delta y_{it} = \theta \Delta x_{it} + \gamma_1 \Delta y_{it-1} + \cdots + \gamma_k \Delta y_{it-k} + \Delta v_{it}. \quad (26)$$

For simplicity and to avoid many instruments problems – see e.g. Roodman (2009) – we focus here on estimation using the instrument sets

$$\mathbf{z}'_{it}(\ell, P) \equiv \begin{bmatrix} y_{it-2} & \cdots & y_{it-(\ell+1)} & x_{it-1} \end{bmatrix} \quad \mathbf{z}'_{it}(\ell, S) \equiv \begin{bmatrix} \mathbf{z}'_{it}(\ell, P) & x_{it} \end{bmatrix} \quad (27)$$

similar to Anderson and Hsiao (1982). Modulo a change in notation, one could just as easily proceed using the instrument sets suggested by Arellano and Bond (1991). Throughout this discussion we use ℓ as a placeholder for the lag length used in estimation. If $\ell = 0$, $\mathbf{z}'_{it}(0, P) = x_{it-1}$ and $\mathbf{z}'_{it}(0, S) = (x_{it-1}, x_{it})$. Given these instrument sets, we have $(\ell + 1) \times (T - \ell - 1)$ moment conditions if x is assumed to be predetermined versus $(\ell + 2) \times (T - \ell - 1)$ if it is assumed to be strictly exogenous, corresponding to the instrument matrices $Z_i(\ell, P) = \text{diag} \{ \mathbf{z}'_{it}(\ell, P) \}_{t=\ell+2}^T$ and $Z_i(\ell, S) = \text{diag} \{ \mathbf{z}'_{it}(\ell, S) \}_{t=\ell+2}^T$. To abstract for a moment from the model selection decision, suppose that we estimate a model with the true lag length: $\ell = k$. The only difference between the P and S sets of moment conditions is that the latter adds over-identifying information in the form of $E[x_{it}\Delta v_{it}]$. If x is strictly exogenous, this expectation equals zero, but if x is only predetermined, then $E[x_{it}\Delta v_{it}] = -E[x_{it}v_{it-1}] \neq 0$ so the over-identifying moment condition is invalid. Given the instrument sets that we consider, this is the only violation of strict exogeneity that is relevant for our moment selection so we take $E[x_{it}v_{it-1}] = -\tau/\sqrt{n}$.

In the examples and simulations described below we consider two-stage least squares (TSLS) estimation of μ_{SR} and μ_{LR} using the instruments defined in Equation 27. For simplicity, we select between two lag length specifications: the first is correct, $\ell = k$, and the second includes one lag too few: $\ell = r$ where $r = k - 1$. Accordingly, we make the coefficient associated with the k^{th} lag local to zero. Let $\boldsymbol{\gamma}' = (\gamma_1, \dots, \gamma_{k-1}, \gamma_k)$ denote the full vector of lag coefficients and $\boldsymbol{\gamma}_r = (\gamma_1, \dots, \gamma_r)$ denote the first $r = k - 1$ lag coefficients. Then, the true parameter vector is $\beta_n = (\theta, \boldsymbol{\gamma}'_r, \delta/\sqrt{n})'$ which becomes, in the limit, $\beta = (\theta, \boldsymbol{\gamma}'_r, 0)'$. To indicate the subvector of β that excludes the k^{th} lag coefficient, let $\beta_r = (\theta, \boldsymbol{\gamma}'_r)'$.

Because the two lag specifications we consider use different time periods in estimation, we require some additional notation to make this clear. First let $\Delta \mathbf{y}_i = [\Delta y_{i,k+2}, \dots, \Delta y_{iT}]'$ and $\Delta \mathbf{y}_i^+ = [\Delta y_{i,k+1}, \Delta y_{i,k+2}, \dots, \Delta y_{iT}]'$ where the superscript “+” indicates the inclusion of an additional time period: $t = k + 1$. Define $\Delta \mathbf{x}_i$, $\Delta \mathbf{x}_i^+$, $\Delta \mathbf{v}_i$, and $\Delta \mathbf{v}_i^+$ analogously. Next, define $L^k \Delta \mathbf{y}_i^+ = [\Delta y_{i1}, \Delta y_{i2}, \dots, \Delta y_{iT-k}]'$ where L^k denotes the element-wise application of the k^{th} order lag operator. Note that the first element of $L^k \Delta \mathbf{y}_i^+$ is unobserved since

$\Delta y_{i1} = y_{i1} - y_{i0}$ but $t = 1$ is the first time period. Now we define the matrices of regressors for the two specifications:

$$\begin{aligned} W_i^{+'}(r) &= \begin{bmatrix} \Delta \mathbf{x}_i^+ & L\Delta \mathbf{y}_i^+ & L^2\Delta \mathbf{y}_i^+ & \cdots & L^{k-1}\Delta \mathbf{y}_i^+ \end{bmatrix} \\ W_i'(k) &= \begin{bmatrix} \Delta \mathbf{x}_i & L\Delta \mathbf{y}_i & L^2\Delta \mathbf{y}_i & \cdots & L^{k-1}\Delta \mathbf{y}_i & L^k\Delta \mathbf{y}_i \end{bmatrix}. \end{aligned}$$

Note that $W_i^{+'}(r)$ contains one more row than $W_i(k)$ but $W_i(k)$ contains one more column than $W_i^{+'}(r)$: removing the k^{th} lag from the model by setting $\ell = r = k - 1$ allows us to use an additional time period in estimation and reduces the number of regressors by one. Stacking over individuals, let $\Delta \mathbf{y} = [\Delta \mathbf{y}'_1 \cdots \Delta \mathbf{y}'_n]'$, $W_\ell = [W_1(\ell) \cdots W_n(\ell)]'$ and define $\Delta \mathbf{y}^+$ and W_ℓ^+ analogously, where ℓ denotes the lag length used in estimation. Finally, let $Z'(\ell, \cdot) = [Z'_1(\ell, \cdot) \cdots Z'_n(\ell, \cdot)]$ where (\cdot) is P or S depending on the instrument set in use. Using this notation, under local mis-specification the true model is

$$\Delta \mathbf{y} = W(k)\beta_n + \Delta \mathbf{v} \quad \Delta \mathbf{y}^+ = W(k)^+\beta_n + \Delta \mathbf{v}^+ \quad (28)$$

Using the shorthand $\widehat{Q} \equiv n[W'Z(Z'Z)^{-1}Z'W]^{-1}W'Z(Z'Z)^{-1}$ our candidate estimators are

$$\widehat{\beta}(k, \cdot) = \widehat{Q}(k, \cdot) \left[\frac{Z'(k, \cdot)\Delta \mathbf{y}}{n} \right] \quad \widehat{\beta}(r, \cdot) = \widehat{Q}(r, \cdot) \left[\frac{Z'(r, \cdot)\Delta \mathbf{y}^+}{n} \right] \quad (29)$$

where (\cdot) is either P or S depending on which instrument set is used and $r = k - 1$, one lag fewer than the true lag length k . The following result describes the limit distribution of $\widehat{\beta}(k, P)$, $\widehat{\beta}(k, S)$, $\widehat{\beta}(r, P)$, and $\widehat{\beta}(r, S)$ which we will use to construct the GFIC.

Theorem 5.2 (Limit Distributions for Dynamic Panel Estimators). *Let $(y_{nit}, x_{nit}, v_{nit})$ be a triangular array of random variables that is iid over i , stationary over t , and satisfies Equation 26 with $\gamma_k = \delta/\sqrt{n}$. Suppose further that x_{it} is predetermined with respect to v_{it} but not strictly exogenous: $E[x_{it}\Delta v_{it}] = \tau/\sqrt{n}$. Then, under standard regularity conditions,*

$$\begin{aligned} \sqrt{n} [\widehat{\beta}(k, P) - \beta] &\rightarrow^d \begin{bmatrix} 0 & \mathbf{0}'_{k-1} & \delta \end{bmatrix}' + Q(k, P) N(\mathbf{0}, \mathcal{V}(k, P)) \\ \sqrt{n} [\widehat{\beta}(k, S) - \beta] &\rightarrow^d \begin{bmatrix} 0 & \mathbf{0}'_{k-1} & \delta \end{bmatrix}' + Q(k, S) \left\{ \boldsymbol{\iota}_{T-(k+1)} \otimes \begin{bmatrix} \mathbf{0}_{k+1} \\ \tau \end{bmatrix} + N(\mathbf{0}, \mathcal{V}(k, S)) \right\} \\ \sqrt{n} [\widehat{\beta}(r, P) - \beta_r] &\rightarrow^d Q(r, P) [\boldsymbol{\iota}_{T-k} \otimes \delta \boldsymbol{\psi}_P + N(\mathbf{0}, \mathcal{V}(r, P))] \\ \sqrt{n} [\widehat{\beta}(r, S) - \beta_r] &\rightarrow^d Q(r, S) \left[\boldsymbol{\iota}_{T-k} \otimes \left(\delta \begin{bmatrix} \boldsymbol{\psi}_P \\ \boldsymbol{\psi}_S \end{bmatrix} + \begin{bmatrix} \mathbf{0}_k \\ \tau \end{bmatrix} \right) + N(\mathbf{0}, \mathcal{V}(r, S)) \right] \end{aligned}$$

where $k = r - 1$, $\beta' = (\theta, \gamma_1, \dots, \gamma_r, 0)$, $\beta'_r = (\theta, \gamma_1, \dots, \gamma_r)$, $\mathcal{V}(k, \cdot) = \text{Var}[Z_i(k, \cdot)\Delta\mathbf{v}_i]$, $\mathcal{V}(r, \cdot) = \text{Var}[Z_i(r, \cdot)\Delta\mathbf{v}_i^+]$, $\widehat{Q}(\ell, \cdot) \rightarrow_p Q(\ell, \cdot)$, $\boldsymbol{\psi}_P = E[\mathbf{z}_{it}(r, P)\Delta y_{it-k}]$, $\psi_S = E[x_{it}\Delta y_{it-k}]$, $\mathbf{z}_{it}(\ell, \cdot)$ is as in Equation 27, $Z_i(\ell, \cdot) = \text{diag}\{\mathbf{z}'_{it}(\ell, \cdot)\}_{t=\ell+2}^T$ and $\boldsymbol{\iota}_d$ denotes a d -vector of ones.

To operationalize the GFIC, we need to provide appropriate estimators of all quantities that appear in Theorem 5.2. To estimate $Q(k, P)$, $Q(k, S)$, $Q(r, P)$, and $Q(r, S)$ we employ the usual sample analogues $\widehat{Q}(\cdot, \cdot)$ given above, which remain consistent under local mis-specification. There are many consistent estimators for the variance matrices $\mathcal{V}(k, P)$, $\mathcal{V}(k, S)$, $\mathcal{V}(r, P)$, $\mathcal{V}(r, S)$ under local mis-specification. In our simulations below, we employ the usual heteroskedasticity-consistent, panel-robust variance matrix estimator. Because $E[\mathbf{z}_{it}(\ell, S)\Delta v_{it}] \neq 0$, we center our estimators of $\mathcal{V}(\ell, S)$ by subtracting the sample analogue of this expectation when calculating the sample variance. We estimate $\boldsymbol{\psi}_P$ and ψ_S as follows

$$\widehat{\boldsymbol{\psi}}_P = \frac{1}{n(T-k-1)} \sum_{t=k+2}^T \sum_{i=1}^n \mathbf{z}_{it}(r, P)\Delta y_{it-k} \quad \widehat{\psi}_S = \frac{1}{n(T-k-1)} \sum_{t=k+2}^T \sum_{i=1}^n x_{it}\Delta y_{it-k}$$

using our assumption of stationarity from above. The only remaining quantities we need to construct the GFIC involve the bias parameters δ and τ . We can read off an asymptotically unbiased estimator of δ directly from Theorem 5.2, namely $\widehat{\delta} = \sqrt{n} \widehat{\gamma}_k(k, P)$ the estimator of γ_k based on the instrument set that assumes only that x is pre-determined rather than strictly exogenous. To construct an asymptotically unbiased estimator of τ , we use the residuals from the specification that uses *both* the correct moment conditions and the correct lag specification, specifically

$$\widehat{\tau} = \left(\frac{\boldsymbol{\iota}'_{T-k-1}}{T-k-1} \right) n^{-1/2} X' \left[\Delta \mathbf{y} - W(k) \widehat{\beta}(k, P) \right] \quad (30)$$

where $X' = [X_1 \cdots X_n]$ and $X_i = \text{diag}\{x_{it}\}_{t=k+2}^T$. The following result gives the joint limiting behavior of $\widehat{\delta}$ and $\widehat{\tau}$, which we will use to construct the GFIC.

Theorem 5.3 (Joint Limit Distribution of $\widehat{\delta}$ and $\widehat{\tau}$). *Under the conditions of Theorem 5.2,*

$$\begin{bmatrix} \widehat{\delta} - \delta \\ \widehat{\tau} - \tau \end{bmatrix} \xrightarrow{d} \Psi N(\mathbf{0}, \Pi \mathcal{V}(k, S) \Pi')$$

where $\widehat{\delta} = \sqrt{n}[\mathbf{e}'_k \widehat{\beta}(k, P)]$, $\mathbf{e}_k = (0, \mathbf{0}'_{k-1}, 1)'$, $\widehat{\tau}$ is as defined in Equation 30,

$$\Psi = \begin{bmatrix} \left(\frac{\boldsymbol{\iota}'_{T-k-1}}{T-k-1} \right) \{ \boldsymbol{\xi}' Q(k, P) \otimes \boldsymbol{\iota}'_{T-k-1} \} & \left(\frac{\boldsymbol{\iota}'_{T-k-1}}{T-k-1} \right) \\ \mathbf{e}'_k Q(k, P) & \mathbf{0}'_{T-k-1} \end{bmatrix},$$

$\xi' = E \left\{ x_{it} \begin{bmatrix} \Delta x_{it} & L\Delta y_{it} & \cdots & L^k \Delta y_{it} \end{bmatrix} \right\}$, the variance matrix $\mathcal{V}(k, S)$ is as defined in Theorem 5.2, the permutation matrix $\Pi = \begin{bmatrix} \Pi'_1 & \Pi'_2 \end{bmatrix}'$ with $\Pi_1 = I_{T-k-1} \otimes \begin{bmatrix} I_{k+1} & \mathbf{0}_{k+1} \end{bmatrix}$ and $\Pi_2 = I_{T-k-1} \otimes \begin{bmatrix} \mathbf{0}'_{k+1} & 1 \end{bmatrix}$, $\boldsymbol{\iota}_d$ is a d -vector of ones and I_d the $(d \times d)$ identity matrix.

To provide asymptotically unbiased estimators of the quantities τ^2 , δ^2 and $\tau\delta$ that appear in the AMSE expressions for our estimators, we apply a bias correction to the asymptotically unbiased estimators of δ and τ from Theorem 5.3.

Corollary 5.1. *Let $\widehat{\Psi}$ be a consistent estimator of Ψ , defined in Theorem 5.3, and $\widehat{\mathcal{V}}(k, S)$ be a consistent estimator of $\mathcal{V}(k, S)$, defined in Theorem 5.2. Then, the elements of*

$$\begin{bmatrix} \widehat{\delta^2} & \widehat{\delta\tau} \\ \widehat{\tau\delta} & \widehat{\tau^2} \end{bmatrix} - \widehat{\Psi} \Pi \widehat{\mathcal{V}}(k, S) \Pi' \widehat{\Psi}'$$

provide asymptotically unbiased estimators of δ^2 , τ^2 and $\tau\delta$, where Π is the permutation matrix defined in Theorem 5.3.

We have already discussed consistent estimation of $\widehat{\mathcal{V}}(k, S)$. Since Π is a known permutation matrix, it remains only to propose a consistent estimator of Ψ . The matrix Ψ , in turn, depends only on $Q(k, P)$, and ξ' . The sample analogue $\widehat{Q}(k, P)$ is a consistent estimator for $Q(k, P)$, as mentioned above, and

$$\widehat{\xi}' = \frac{1}{n(T-k-1)} \sum_{t=k+2}^T \sum_{i=1}^n x_{it} \begin{bmatrix} \Delta x_{it} & L\Delta y_{it} & \cdots & L^k \Delta y_{it} \end{bmatrix} \quad (31)$$

is consistent for ξ' . We now have all the quantities needed to construct the GFIC for μ_{SR} , the short-run effect of x on y . Since $\mu_{SR} = \theta$, we can read off the AMSE expression for this parameter directly from Theorem 5.2. For the long-run effect μ_{LR} , however, we need to formally apply the Delta-method and account for the fact that the true value of γ_k is δ/\sqrt{n} . Expressed as a function φ of the underlying model parameters,

$$\mu_{LR} = \varphi(\theta, \boldsymbol{\gamma}_r, \gamma_k) = \frac{\theta}{1 - \boldsymbol{\iota}'_r \boldsymbol{\gamma}_r - \gamma_k}$$

and the derivatives of φ are

$$\nabla \varphi(\theta, \boldsymbol{\gamma}_r, \gamma_k)' = \begin{bmatrix} \frac{\partial \varphi}{\partial \theta} & \frac{\partial \varphi}{\partial \boldsymbol{\gamma}'_r} & \frac{\partial \varphi}{\partial \gamma_k} \end{bmatrix} = \left(\frac{1}{1 - \boldsymbol{\iota}'_r \boldsymbol{\gamma}_r - \gamma_k} \right)^2 \begin{bmatrix} (1 - \boldsymbol{\iota}'_r \boldsymbol{\gamma}_r - \gamma_k) & \theta \boldsymbol{\iota}'_r & \theta \end{bmatrix}.$$

Using this notation, the limiting value of μ_{LR} is $\mu_{LR}^0 = \varphi(\theta, \boldsymbol{\gamma}_r, 0)$ while the true value is $\mu_{LR}^n = \varphi(\theta, \boldsymbol{\gamma}_r, \delta/\sqrt{n})$. Similarly, the limiting value of $\nabla \varphi$ is $\nabla \varphi_0 = \nabla \varphi(\theta, \boldsymbol{\gamma}_r, 0)'$, obtained

by putting zero in place of γ_k . We estimate this quantity consistently by plugging in the estimates from $\hat{\beta}(k, P)$.

5.3 Slope Heterogeneity Example

Suppose we wish to estimate the average effect β of a regressor x in a panel setting where this effect may vary by individual: say $\beta_i \sim \text{iid}$ over i . One idea is to simply ignore the heterogeneity and fit a pooled model. A pooled estimator will generally be quite precise, but depending on the nature and extent of heterogeneity could show a serious bias. Another idea is to apply the *mean group estimator* by running separate time-series regressions for each individual and averaging the result over the cross-section (Pesaran et al., 1999; Pesaran and Smith, 1995; Swamy, 1970). This approach is robust to heterogeneity but may yield an imprecise estimator, particularly in panels with a short time dimension. To see how the GFIC navigates this tradeoff, consider the following DGP:

$$y_{it} = \beta_i x_{it} + \epsilon_{it} \quad (32)$$

$$\beta_i = \beta + \eta_i, \quad \eta_i \sim \text{iid}(0, \sigma_\eta^2) \quad (33)$$

where x_{it} is uncorrelated with ϵ_{it} but is *not* assumed to be independent of η_i . As in the preceding examples $i = 1, \dots, n$ indexes individuals, $t = 1, \dots, T$ indexes time periods, and we assume without loss of generality that all random variables are mean zero and any exogenous controls have been projected out. For the purposes of this example, assume further that ϵ_{it} is iid over both i and t with variance σ_ϵ^2 and that both error terms are homoskedastic: $E[\epsilon_{it}^2 | x_{it}] = \sigma_\epsilon^2$ and $E[\eta_i^2 | x_{it}] = \sigma_\eta^2$, and $E[\eta_i \epsilon_{it} | x_{it}] = 0$. Neither homoskedasticity nor time-independent errors are required to apply the GFIC to this example, but these assumptions simplify the exposition. We place no assumptions on the joint distribution of x_{it} and η_i .

Stacking observations, let $\mathbf{y}_i = (y_{i1}, \dots, y_{iT})'$ and define \mathbf{x}_i analogously. We consider two estimators: the pooled OLS estimator $\hat{\beta}_{OLS}$ and the mean-group estimator $\hat{\beta}_{MG}$

$$\hat{\beta}_{OLS} = \left(\sum_{i=1}^n \mathbf{x}_i' \mathbf{x}_i \right)^{-1} \left(\sum_{i=1}^n \mathbf{x}_i' \mathbf{y}_i \right) \quad (34)$$

$$\hat{\beta}_{MG} = \frac{1}{n} \sum_{i=1}^n \hat{\beta}_i = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i' \mathbf{x}_i)^{-1} (\mathbf{x}_i' \mathbf{y}_i) \quad (35)$$

where $\hat{\beta}_i$ denotes the OLS estimator calculated using observations for individual i only. If we knew with certainty that there was no slope heterogeneity, we could clearly prefer $\hat{\beta}_{OLS}$ as it is both unbiased and has the lower variance. In the presence of heterogeneity, however,

the situation is more complicated. If $E[\mathbf{x}'_i \mathbf{x}_i \eta_i] \neq 0$ then $\hat{\beta}_{OLS}$ will show a bias whereas $\hat{\beta}_{MG}$ will not. To encode this idea within the local mis-specification framework, we take $E[\mathbf{x}'_i \mathbf{x}_i \eta_i] = \tau/\sqrt{n}$ so that, for any fixed n the OLS estimator is biased unless $\tau = 0$ but this bias disappears in the limit. Turning our attention from bias to variance, we might expect that $\hat{\beta}_{OLS}$ would remain the more precise estimator in the presence of heterogeneity. In fact, however, this need not be the case: as we show below, $\hat{\beta}_{MG}$ will have a *lower* variance than $\hat{\beta}_{OLS}$ if σ_η^2 is sufficiently large. To construct the GFIC for this example, we estimate the bias parameter τ by substituting the mean group estimator into the OLS moment condition:

$$\hat{\tau} = \frac{1}{\sqrt{n}} \sum_{i=1}^n \mathbf{x}'_i (\mathbf{y}_i - \mathbf{x}_i \hat{\beta}_{MG}). \quad (36)$$

The key result needed to apply the GFIC in this this example gives the joint limiting distribution of $\hat{\tau}$, the mean-group estimator, and the OLS estimator.

Theorem 5.4 (Limit Distribution of OLS and Mean-Group Estimators). *Let $(\mathbf{x}_{ni}, \eta_{ni}, \epsilon_{ni})$ be an iid triangular array of random variables such that $E[\mathbf{x}'_{ni} \epsilon_{ni}] = 0$, $Var(\epsilon_{ni} | \mathbf{x}_{ni}) \rightarrow \sigma_\epsilon^2 I_T$, $Var(\eta_{ni} | \mathbf{x}_{ni}) \rightarrow \sigma_\eta^2$, $E[\eta_{ni} \epsilon_{ni} | \mathbf{x}_{ni}] \rightarrow 0$, and $E[\mathbf{x}'_{ni} \mathbf{x}_{ni} \eta_{ni}] = \tau/\sqrt{n}$. Then, under standard regularity conditions,*

$$\begin{bmatrix} \sqrt{n}(\hat{\beta}_{OLS} - \beta) \\ \sqrt{n}(\hat{\beta}_{MG} - \beta) \\ \hat{\tau} \end{bmatrix} \xrightarrow{d} N \left(\begin{bmatrix} \tau/\kappa \\ 0 \\ \tau \end{bmatrix}, \begin{bmatrix} \left(\frac{\lambda^2 + \kappa^2}{\kappa^2} \right) \sigma_\eta^2 + \frac{\sigma_\epsilon^2}{\kappa} & \sigma_\eta^2 + \frac{\sigma_\epsilon^2}{\kappa} & \left(\frac{\lambda^2}{\kappa} \right) \sigma_\eta^2 \\ \sigma_\eta^2 + \zeta \sigma_\epsilon^2 & \sigma_\eta^2 + \zeta \sigma_\epsilon^2 & \sigma_\epsilon^2 (1 - \kappa \zeta) \\ \lambda^2 \sigma_\eta^2 + \kappa(\kappa \zeta - 1) \sigma_\epsilon^2 & \sigma_\epsilon^2 (1 - \kappa \zeta) & \lambda^2 \sigma_\eta^2 + \kappa(\kappa \zeta - 1) \sigma_\epsilon^2 \end{bmatrix} \right)$$

where $\kappa = E[\mathbf{x}'_i \mathbf{x}_i]$, $\lambda^2 = Var(\mathbf{x}'_i \mathbf{x}_i)$, and $\zeta = E[(\mathbf{x}'_i \mathbf{x}_i)^{-1}]$.

As mentioned above, the OLS estimator need not have a lower variance than the mean-group estimator if σ_η^2 is sufficiently large. This fact follows as a corollary of Theorem 5.4.

Corollary 5.2. *Under the conditions of Theorem 5.4, the asymptotic variance of the OLS estimator is lower than that of the mean-group estimator if and only if $\lambda^2 \sigma_\eta^2 < \sigma_\epsilon^2 (\kappa^2 \zeta - \kappa)$, where $\kappa = E[\mathbf{x}'_i \mathbf{x}_i]$, $\lambda^2 = Var(\mathbf{x}'_i \mathbf{x}_i)$, and $\zeta = E[(\mathbf{x}'_i \mathbf{x}_i)^{-1}]$.*

Note, as a special case of the preceding, that the OLS estimator is guaranteed to have the lower asymptotic variance when $\sigma_\eta^2 = 0$ since $E[\mathbf{x}'_i \mathbf{x}_i]^{-1} < E[(\mathbf{x}'_i \mathbf{x}_i)^{-1}]$ by Jensen's inequality. When $\sigma_\eta^2 \neq 0$, the situation is in general much more complicated. A simple normal example, however, provides some helpful intuition. Suppose that for a given individual the observations x_{it} are iid standard normal over t . Then $\mathbf{x}'_i \mathbf{x}_i \sim \chi_T^2$, so that $\kappa = T$, $\lambda^2 = 2T$ and $\zeta = 1/(T-2)$, provided of course that $T > 2$. Substituting these into Corollary 5.2, the OLS estimator will

have the lower asymptotic variance whenever $(T - 2)\sigma_\eta^2 < \sigma_\varepsilon^2$. All else equal, the shorter the panel, the more likely that OLS will have the lower variance. But if σ_η^2 is large enough, the length of the panel becomes irrelevant: with enough slope heterogeneity, the mean-group estimator has the advantage both in bias and variance.

To apply the GFIC in practice, we first need to determine whether the OLS estimator has the smaller asymptotic variance. This requires us to estimate the quantities λ^2 , κ , and ζ from Theorem 5.4 along with σ_η^2 and σ_ε^2 . The following estimators are consistent under the assumptions of Theorem 5.4:

$$\begin{aligned}\hat{\kappa} &= \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i' \mathbf{x}_i & \hat{\zeta} &= \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i' \mathbf{x}_i)^{-1} \\ \hat{\lambda}^2 &= \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i' \mathbf{x}_i - \hat{\kappa})^2 & \hat{\sigma}_\varepsilon^2 &= \frac{1}{nT-1} \sum_{i=1}^n \sum_{t=1}^T (y_{it} - x_{it} \hat{\beta}_{OLS})^2 \\ \hat{\sigma}_\eta^2 &= \frac{S_b}{n-1} - \frac{1}{n} \sum_{i=1}^n \hat{\sigma}_\varepsilon^2 (\mathbf{x}_i' \mathbf{x}_i)^{-1} & S_b &= \sum_{i=1}^n \hat{\beta}_i^2 - n \left(\frac{1}{n} \sum_{i=1}^n \hat{\beta}_i \right)^2\end{aligned}$$

If the estimated asymptotic variance of the mean-group estimator is lower than that of the OLS estimator, then there is no need to estimate AMSE: we should simply use the mean-group estimator. If this is not the case, then we construct the GFIC using the asymptotically unbiased estimator $\hat{\tau}^2 - \hat{\sigma}_\tau^2$ of τ^2 , where $\hat{\sigma}_\tau^2 = \hat{\lambda}^2 \hat{\sigma}_\eta^2 + \hat{\kappa}(\hat{\kappa} \hat{\zeta} - 1) \hat{\sigma}_\varepsilon^2$ is a consistent estimator of the asymptotic variance of $\hat{\tau}$.

6 Simulation Results

We now evaluate the performance of the GFIC in a number of simulation experiments based on the Examples from Section 5.⁸

6.1 Fixed vs. Random Effects Example

We employ a simulation design similar to that used by Guggenberger (2010), namely

$$y_{it} = 0.5x_{it} + \alpha_i + \varepsilon_{it}$$

⁸In the interest of space we omit the simulation results for the slope heterogeneity example from Section 5.3. These are available upon request.

where

$$\begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{iT} \\ \alpha_i \end{bmatrix} \stackrel{\text{iid}}{\sim} N \left(\begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & \rho & \dots & \rho & \gamma \\ \rho & 1 & \dots & \rho & \gamma \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ \rho & \rho & \dots & 1 & \gamma \\ \gamma & \gamma & \dots & \gamma & 1 \end{bmatrix} \right)$$

independently of $(\varepsilon_{i1}, \dots, \varepsilon_{iT})' \sim \text{iid } N(0, \sigma_\varepsilon^2 \mathbf{I}_T)$. In this design, γ controls the correlation between x_{it} and the individual effects α_i , while ρ controls the persistence of x_{it} over time. Larger values of γ correspond to larger violations of the assumption underlying the random effects estimator, increasing its bias. Larger values of ρ , on the other hand, decrease the amount of variation within individuals, thus *increasing* the variance of the fixed effects estimator. Figure 1 presents RMSE values for the random effects GLS estimator and fixed effects estimator along with those for the post-GFIC and averaging estimators described above in Section 5.1 over a grid of values for γ , ρ and n with $T = 2$. Results for $T = 5$ appear in Figure 4 of Appendix B. All calculations are based on 10,000 simulation replications. In the interest of space, we present only results for $\sigma_\varepsilon^2 = 2.5$ and a coarse parameter grid for ρ here. Additional results are available upon request.

We see from Figure 1 that, regardless of the configuration of the other parameter values, there is always a range of values for γ for which the random effects estimator has a smaller RMSE than the fixed effects estimator. The width of this range increases as either the number of individuals N or the number of time periods T decrease. It also increases as the persistence ρ of x_{it} increases. Indeed, when N and T are relatively small and ρ is relatively large, the individual effects α_i can be *strongly* correlated with x_{it} and still result in a random effects estimator with a lower RMSE than the fixed effects estimator. The post-GFIC estimator essentially “splits the difference” between the random and fixed effects estimators. While it cannot provide a uniform improvement over the fixed effects estimator, the post-GFIC estimator performs well. When γ is not too large it can yield a substantially lower RMSE than the fixed effects estimator. The gains are particularly substantial when x_{it} is relatively persistent and T relatively small, as is common in micro-panel datasets. The averaging estimator performs even better, providing a nearly uniform improvement over the post-GFIC estimator. Only at very large values of γ does it yield a higher RMSE, and these are points in the parameter space where the fixed effects, post-GFIC and averaging estimators are for all intents and purposes identical in RMSE. Results for $T = 5$ are qualitatively similar. See Figure 4 of Appendix B for details.




Figure 1: RMSE values for the Random vs. Fixed effects simulation example from Section 6.1 with $T = 2$ and $\sigma_\varepsilon^2 = 2.5$. Results are based on 10,000 simulation replications.

6.2 Dynamic Panel Example

We now consider two simulation experiments based on section 5.2, applying the GFIC to a dynamic panel model. For both experiments our data generating process is similar to that of Andrews and Lu (2001), specifically

$$\begin{bmatrix} x_i \\ \eta_i \\ v_i \end{bmatrix} \sim \text{iid } N \left(\begin{bmatrix} \mathbf{0}_T \\ 0 \\ \mathbf{0}_T \end{bmatrix}, \begin{bmatrix} I_T & \sigma_{x\eta}\iota_T & \sigma_{xv}\Gamma_T \\ \sigma_{x\eta}\iota_T' & 1 & \mathbf{0}_T' \\ \sigma_{xv}\Gamma_T' & \mathbf{0}_T & I_T \end{bmatrix} \right) \quad (37)$$

where $\mathbf{0}_m$ denotes an m -vector of zeros, I_m the $(m \times m)$ identity matrix, ι_m an m -vector of ones, and Γ_m an $m \times m$ matrix with ones on the subdiagonal and zeros elsewhere, namely

$$\Gamma_m = \begin{bmatrix} \mathbf{0}_{m-1}' & 0 \\ I_{m-1} & \mathbf{0}_{m-1} \end{bmatrix}. \quad (38)$$

Under this covariance matrix structure η_i and v_i are uncorrelated with each other, but both are correlated with x_i : $E[x_{it}\eta_i] = \sigma_{x\eta}$ and x_{it} is predetermined but not strictly exogenous with respect to v_{it} . Specifically, $E[x_{it}v_{it-1}] = \sigma_{xv}$, while $E[x_{it}v_{is}] = 0$ for $s \neq t-1$. We initialize the pre-sample observations of y to zero, the mean of their stationary distribution, and generate the remaining time periods according to Equation 24 with $\theta = 0.5$ and $\sigma_{x\eta} = 0.2$. The true lag length differs in our two examples as does the target parameter, so we explain these features of the simulation designs below. Unlike Andrews and Lu (2001) we do not generate extra observations to keep the time dimension fixed across estimators with different lag specifications. This is for two reasons. First, in real-world applications such additional observations would not be available. Second, we are explicitly interested in trading off the efficiency gain from including additional time periods in estimation against the bias that arises from estimating an incorrect lag specification.

6.2.1 Long-run versus Short-run Effects

Consider two different researchers who happen to be working with the same panel dataset. One wishes to estimate the short-run effect of x on y while the other wishes to estimate the long-run effect. Should they use the same model specification? We now present an example showing that the answer, in general, is no. Suppose that the true model is

$$y_{it} = \theta x_{it} + \gamma_1 y_{it-1} + \gamma_2 y_{it-2} + \eta_i + v_{it}$$

where $i = 1, \dots, n = 250$ and $t = 1, \dots, T = 5$ and the regressor, individual effect and error term are generated according to Equation 37, as described in the preceding section. Our model selection decision in this example is whether to set $\gamma_2 = 0$ and estimate a specification with one lag only. We denote this one-lag specification by L1 and the true specification, including both lags, by L2. To focus on the model selection decision, we fix the instrument set in this experiment to $\mathbf{z}_{it}(\ell, P)$, defined in Equation 27. Because this instrument set is valid when x is pre-determined, it does not introduce bias into our estimation. Thus, bias only emerges if we estimate L1 when $\gamma_2 \neq 0$. Our simulation design takes $\theta = 0.5$, $\gamma_1 = 0.4$, $\sigma_{x\eta} = 0.2$, and $\sigma_{xv} = 0.1$ and varies γ_2 over the range $\{0.10, 0.11, \dots, 0.19, 0.20\}$.

Table 1 presents the results of the simulation, based on 1000 replications at each grid point. Because they are based on *ratios* of estimators of θ and γ_1, γ_2 , estimators of the long-run effect may not have finite moments, making finite-sample MSE undefined. The usual solution to this problem in simulation settings is to work with so-called “trimmed” MSE by discarding observations that fall outside, say, a range $[-M, M]$ before calculating MSE.⁹ Because there is no clear way to set the trimming constant M , it can be difficult to interpret results based on trimmed MSE unless one considers a variety of values of M . To avoid this issue, Table 1 reports simulation results for median absolute deviation (MAD). Results for trimmed MSE with different choices of M are similar and are available upon request.

The columns of Table 1 labeled L1 and L2 give the MAD of estimators that fix the lag length to one and two, while those labeled GFIC give the MAD of an estimator that selects lag length via the GFIC. Notice that throughout the table $\gamma_2 \neq 0$ so that L1 is *mis-specified*. Nonetheless, L1 yields lower MAD estimators of both the short-run and long-run effects when γ_2 is sufficiently small and the difference can be substantial. When $\gamma_2 = 0.2$, for example, MAD for is 0.582 for the long-run effect estimator based on L1 versus 0.801 for that based on L2. Note moreover that the point at which γ_2 becomes large enough for L2 to be preferred depends on which effect we seek to estimate. When γ_2 equals 0.15 or 0.16, L1 gives a lower MAD for the short-run effect while L2 gives a lower MAD for the long-run effect. Because it is subject to random model selection errors, the GFIC can never outperform the oracle estimator that uses L1 when it is optimal in terms of MAD and L2 otherwise. Instead, the GFIC represents a compromise between two extremes: its MAD is never as large as that of the worst specification and never as small as that of the best specification. When there are large MAD differences between L1 and L2, however, GFIC is generally quite close to the

⁹Note that *asymptotic* MSE remains well-defined even for estimators that do not possess finite-sample moments so that GFIC comparisons remain meaningful. By taking the trimming constant M to infinity, one can formalize the notion that asymptotic MSE comparisons can be used to “stand in” for finite-sample MSE even when the latter does not exist. For more details, See Hansen (2016) and online appendix C of DiTraglia (2016).

γ_2	Short-run Effect			Long-run Effect		
	L2	L1	GFIC	L2	L1	GFIC
0.10	0.231	0.141	0.173	0.801	0.582	0.688
0.11	0.237	0.156	0.181	0.834	0.633	0.716
0.12	0.240	0.174	0.193	0.850	0.685	0.752
0.13	0.238	0.187	0.201	0.870	0.729	0.787
0.14	0.220	0.198	0.203	0.870	0.764	0.808
0.15	0.201	0.219	0.211	0.844	0.822	0.839
0.16	0.205	0.223	0.210	0.883	0.856	0.862
0.17	0.181	0.242	0.204	0.860	0.911	0.897
0.18	0.162	0.258	0.189	0.835	0.959	0.891
0.19	0.161	0.265	0.181	0.866	0.997	0.917
0.20	0.143	0.288	0.162	0.858	1.054	0.910

Table 1: Comparisons of mean absolute deviation (MAD) for estimators of the Short-run and Long-run effects of x on y in the simulation experiment described in Section 6.2.1. The columns labeled L1 and L2 give the MAD of estimators that fix the lag length to one and two, while the columns labeled GFIC give the MAD of an estimator that selects lag length via the GFIC. Results are based on 1000 simulation replications from the DGP described in Section 6.2 with $\gamma_1 = 0.4$, using the estimators described in Section 5.2 and the instrument set $\mathbf{z}_{it}(\ell, P)$ from Equation 27.

optimum.

6.2.2 Simultaneous Model and Moment Selection for the Short-run Effect

We now consider a more complicated simulation experiment that simultaneously selects over lag specification and endogeneity assumptions. In this simulation our target parameter is the short-run effect of x on y , as in our empirical example below and the true model contains one lag. Specifically,

$$y_{it} = \theta x_{it} + \gamma y_{it-1} + \eta_i + v_{it}$$

where $i = 1, \dots, n$ and $t = 1, \dots, T$ and the regressor, individual effect and error term are generated according to Equation 37. Our model selection decision example is whether to set $\gamma = 0$ and estimate a specification without the lagged dependent variable, while our moment selection decision is whether to use only the instrument set $\mathbf{z}_{it}(\ell, P)$, which assumes that x is predetermined, or the instrument set $\mathbf{z}_{it}(\ell, S)$ which assumes that it is strictly exogenous. Both instrument sets are defined in Equation 27. We consider four specifications, each estimated by TSLS using the expressions from section 5.2. The correct specification, LP, estimates both γ and θ using only the “predetermined” instrument set. In contrast, LS estimates both parameters using the “strict exogeneity” instrument set. The specifications P and S set $\gamma = 0$ and estimate only θ , using the predetermined and strictly exogenous instrument sets, respectively. Our simulation design sets $\theta = 0.5$, $\sigma_{x\eta} = 0.2$ and varies

γ , σ_{xv} , T and n over a grid. Specifically, we take $\gamma, \sigma_{xv} \in \{0, 0.005, 0.01, \dots, 0.195, 0.2\}$, $n \in \{250, 500\}$, $T \in \{4, 5\}$.¹⁰ All values are computed based on 2000 simulation replications.

The first question is how the finite sample RMSE of the four *fixed* specifications varies with γ , σ_{xv} , T and n . Figure 2 colors each region of the parameter space according to which of the estimators of θ – LP, LS, P or S – yields the lowest finite-sample RMSE. The saturation of a color indicates the relative difference in RMSE of the best estimator at that point measured against the second-best estimator: darker indicates a larger advantage; lighter indicates a smaller advantage. From the figure and table we see that there are potentially large gains to be had by intentionally using a mis-specified estimator. Indeed, the correct specification LP, depicted in orange, is only optimal when both ρ_{xv} and γ are fairly large relative to sample size. When $T = 4$ and $n = 250$, for example, γ and ρ_{xv} must both exceed 0.10 before LP has the lowest RMSE. Moreover, the advantage of the mis-specified estimators can be substantial, as seen from the top panel of Table 2, which presents RMSE values multiplied by 1000 for ease of reading.¹¹

In practice, of course, we do not know the values of γ , ρ , θ , or the other parameters of the DGP so this comparison of finite-sample RMSE values is infeasible. Instead, we consider using the GFIC to select between the four specifications. Clearly there are gains to be had from estimating a mis-specified model in certain situations. The question remains, can the GFIC identify them? Because it is an efficient rather than consistent selection criterion, the GFIC remains random, even in the limit. This means that the GFIC can never outperform the “oracle” estimator that uses whichever specification gives the lowest finite-sample RMSE in figure 2. Moreover, because our target parameter is a scalar, Stein-type results do not apply: the post-GFIC estimator cannot provide a uniform improvement over the true specification LP. Nevertheless, the post-GFIC estimator can provide a substantial improvement over LP when ρ_{xv} and γ are relatively small relative to sample size, as shown in Figure 3 and in the two leftmost panes of the top panel in Table 2. This is precisely the situation for which the GFIC is intended: a setting in which we have reason to suspect that mis-specification is fairly mild. Moreover, in situations where LP has a substantially lower RMSE than the other estimators, the post GFIC-estimator’s performance is comparable.

To provide a basis for comparison, we now consider results for a number of alternative selection procedures. The first is a “Downward J-test,” which is intended to approximate what applied researchers may do in practice when faced with a model and moment selection decision such as this one. The Downward J-test selects the *most restrictive* specification

¹⁰Setting T no smaller than 4 ensures that MSE exists for all four estimators: the finite sample moments of the TSLS estimator only exist up to the order of over-identification.

¹¹In the interest of space Table 2 uses a coarser simulation grid than Figure 2. Tables of RMSE values over the full simulation grid are available upon request.



Figure 2: Minimum RMSE specification at each combination of parameter values for the simulation experiment from Section 6.2.2. Color saturation at a given grid point indicates RMSE relative to second best specification.

		GFIC												LP				LS				P				S			
				ρ_{xv}		γ		0	0.05	0.1	0.15	0	0.05	0.1	0.15	0	0.05	0.1	0.15	0	0.05	0.1	0.15	0	0.05	0.1	0.15		
4	250	N	T	0	57	60	62	67	72	70	69	68	51	58	74	95	51	53	52	52	42	49	65	86					
				0.05	60	62	67	69	74	71	70	52	59	77	95	58	57	58	56	43	54	72	91						
				0.1	66	67	74	75	77	72	74	71	52	57	76	97	76	73	71	70	47	60	77	97					
				0.15	71	75	78	82	81	77	74	75	53	62	76	99	99	96	92	89	55	69	84	103					
5	250	N	T	0	40	43	48	48	51	49	49	48	37	47	67	90	37	37	36	36	30	40	59	83					
				0.05	43	47	51	49	52	50	51	49	36	47	67	90	45	45	44	43	31	47	65	87					
				0.1	45	53	56	54	52	51	50	50	36	47	69	90	67	64	62	59	37	53	73	92					
				0.15	50	56	59	56	56	53	52	50	36	48	69	92	92	90	85	83	45	63	80	100					
5	250	N	T	0	45	48	54	56	56	56	55	54	42	51	70	91	44	45	44	45	36	44	62	83					
				0.05	48	52	56	55	58	56	55	54	43	52	70	92	52	51	51	48	38	50	68	89					
				0.1	51	57	61	58	59	57	55	55	44	53	72	94	68	66	65	62	42	57	75	95					
				0.15	55	61	65	64	60	60	58	57	44	52	74	94	94	89	85	81	51	64	83	100					
500	250	N	T	0	33	36	40	38	41	40	38	38	31	42	63	86	32	31	32	32	27	36	56	79					
				0.05	35	40	40	38	41	39	38	38	31	42	63	87	42	40	38	38	29	43	62	85					
				0.1	37	44	44	41	42	40	39	40	31	43	66	88	63	60	58	55	35	52	72	91					
				0.15	38	45	44	42	42	42	39	40	31	44	67	90	88	85	80	76	44	62	80	98					

T	N	γ	ρ_{xv}	J-test 5%						J-test 10%						GMM-BIC						GMM-AIC						GMM-HQ							
				0		0.05		0.1		0.15		0		0.05		0.1		0.15		0		0.05		0.1		0.15		0		0.05		0.1		0.15	
4	250	0		44	52	66	80	46	53	68	77	44	51	68	89	52	58	72	81	46	55	70	88												
		0.05		47	55	72	88	49	56	73	86	46	55	73	93	53	60	75	88	48	56	74	93												
		0.1		55	61	78	95	58	63	78	94	50	60	78	98	59	64	79	92	53	60	78	97												
		0.15		68	74	86	102	73	75	86	101	57	69	83	103	66	73	84	100	59	69	84	102												
	500	0		32	41	55	61	34	43	55	57	31	42	61	85	37	47	58	61	34	44	63	79												
		0.05		34	47	64	74	36	48	63	69	32	47	66	88	39	49	64	67	34	47	66	82												
		0.1		46	55	71	86	49	56	71	82	39	53	73	92	44	54	70	78	40	53	73	89												
		0.15		68	67	81	97	73	69	81	95	43	62	80	100	52	61	78	92	43	60	79	98												
5	250	0		37	45	62	74	38	46	61	71	40	49	68	89	43	50	66	75	40	50	69	86												
		0.05		41	51	67	82	42	52	66	78	42	52	70	92	44	54	68	79	42	52	70	89												
		0.1		47	58	74	91	50	59	74	88	44	56	74	95	48	58	73	87	45	56	74	94												
		0.15		63	67	82	97	67	68	82	96	46	58	80	99	54	61	79	93	47	58	80	98												
	500	0		28	37	51	51	29	38	49	46	30	41	62	84	31	41	53	50	30	41	60	73												
		0.05		31	44	59	66	32	44	57	61	31	44	64	86	33	44	57	56	31	44	63	77												
		0.1		41	52	70	83	44	53	69	78	32	48	71	90	37	49	66	70	33	48	70	86												
		0.15		65	65	79	94	70	66	78	90	33	52	76	96	41	55	73	84	33	53	76	94												

Table 2: RMSE values multiplied by 1000 for the simulation experiment from Section 6.2.2.




Figure 3: RMSE of the post-GFIC estimator relative to that of the true specification (LP) in the dynamic panel simulation experiment from Section 6.2.2.

that is not rejected by a J-test test with significance level $\alpha \in \{0.05, 0.1\}$. We test the specifications $\{S, P, LS, LP\}$ *in order* and report the first that is *not rejected*. This means that we only report LP if all the other specifications have been rejected. This procedure is, of course, somewhat *ad hoc* because the significance threshold α is chosen arbitrarily rather than with a view towards some kind of selection optimality. We also consider the GMM model and moment selection criteria of [Andrews and Lu \(2001\)](#):

$$\text{GMM-BIC} = J - (|c| - |b|) \log n$$

$$\text{GMM-AIC} = J - 2(|c| - |b|)$$

$$\text{GMM-HQ} = J - 2.01(|c| - |b|) \log \log n$$

where $|b|$ is the number of parameters estimated, and $|c|$ the number of moment conditions used. Under certain assumptions, it can be shown that both the GMM-BIC and GMM-HQ are consistent: they select the maximal correctly specified estimator with probability approaching one in the limit. To implement these criteria, we calculate the J-test based on the optimal, two-step GMM estimator with a panel robust, heteroscedasticity-consistent, centered covariance matrix estimator for each specification. To compare selection procedures we use the same simulation grid as above, namely $\gamma, \sigma_{xv} \in \{0, 0.005, 0.01, \dots, 0.195, 0.20\}$. Again, each point on the simulation grid is calculated from 2000 simulation replications. The bottom panel of Table 2 presents results for these alternative selection procedures. Figures comparing the GFIC against each alternative in turn appear in Appendix B. There is no clear winner in point-wise RMSE comparisons between the GFIC and its competitors. A substantial difference between the GFIC and its competitors emerges, however, when we examine worst-case RMSE. Here the GFIC clearly dominates, providing the lowest worst-case RMSE across all configurations of T and n . The differences are particularly stark for larger sample sizes. For example, when $T = 5$ and $n = 500$ the worst-case RMSE of GFIC is approximately half that of its nearest competitor: GMM-AIC. The consistent criteria, GMM-BIC and GMM-HQ, perform particularly poorly in terms of worst-case RMSE. This is unsurprising given that the worst-case risk of any consistent selection criteria diverges as sample size increases.¹²

7 Empirical Example: The Demand for Cigarettes

We now consider an empirical example illustrating the GFIC in the dynamic panel setting introduced in Section 5.2. Our exercise is based on [Baltagi et al. \(2000\)](#) who study the

¹²See, e.g., [Leeb and Pötscher \(2008\)](#).

demand for cigarettes using panel data for 46 U.S. states between 1963 and 1992. Their model is

$$\ln C_{it} = \gamma \ln C_{i,t-1} + \theta \ln P_{it} + \beta_1 \ln Y_{it} + \beta_2 \ln Pn_{it} + \eta_i + \lambda_t + v_{it}$$

where C_{it} is the number of packs of cigarettes sold per person aged 16 and above, P_{it} is the real average retail price of a pack of cigarettes, Y_{it} is per capita disposable income, Pn_{it} is the minimum average price of a pack of cigarettes in any state that neighbors state i , η_i is a state fixed effect, and λ_t is a time fixed effect. The lagged dependent variable in this model is meant to capture habit-persistence in cigarette consumption but it is the price elasticity not the habit-persistence *per se* that is of primary interest.

Baltagi et al. (2000) consider an exhaustive list of possible estimators for two target parameters, the short-run price elasticity θ and the long-run price elasticity $\theta/(1 - \gamma)$, and explore how the resulting estimates vary. Here we consider selecting between four alternative estimators of the short-run price elasticity θ , as in the second simulation experiment from Section 6.2. Each specification is estimated by TSLS in first differences, using the expressions from section 5.2.¹³ For simplicity, we assume that the controls $\ln Y_{it}$ and $\ln Pn_{it}$ are exogenous with respect to v_{it} . We focus on two questions. First: what exogeneity assumption should we impose on $\ln P_{it}$? Second: should we allow for habit persistence by estimating γ ?

Our baseline specification, LP, estimates both γ and θ and assumes only that $\ln P_{it}$ is predetermined with respect to v_{it} . This estimator uses the instrument set $\mathbf{z}_{it}(\ell, P)$ with $\ell = 1$ from Equation 27. For the purposes of this exercise we assume that LP is correctly specified. The specification LS also estimates γ , but uses the expanded instrument set $\mathbf{z}_{it}(\ell, S)$ with $\ell = 1$ from equation 27. The additional instruments used in LS are only valid if $\ln P_{it}$ is strictly exogenous with respect to v_{it} . Like LP and LS, the specifications P and S differ in whether or not they impose that $\ln P_{it}$ is predetermined or strictly exogenous. In contrast, however, they set $\gamma = 0$ and estimate a model with no habit persistence ($\ell = 0$). This increases the number of time periods available for estimation.

Estimates and GFIC results for all specifications appear in Table 3. In Panel 3a we use data from 1975–1980 only ($T = 6$). After first-differencing, this leaves 4 time periods for estimation in specifications that include a lag (LP and LS) versus 5 for those that do not (P and S). In panel 3b we use data from 1975–1985 ($T = 11$). After first-differencing this leaves 9 time periods for estimation without a lag versus 10 for estimation with a lag. We choose to artificially shorten the time dimension of the panel from Baltagi et al. (2000) to illustrate

¹³Our baseline specification is similar to the estimator that Baltagi et al. (2000) refer to as FD2SLS. There are two differences however. First, whereas they use lags of the exogenous controls $\ln Y_{it}$ and $\ln Pn_{it}$, our instrument set follows Anderson and Hsiao (1982). Second, whereas Baltagi et al. (2000) appear to have inadvertently omitted the time dummies from their FD2SLS specification, we include them.

(a) 1975–1980 ($T = 6$)					(b) 1975–1985 ($T = 11$)				
	LP	LS	P	S		LP	LS	P	S
$\hat{\theta}$	-0.68	-0.32	-0.28	-0.37	$\hat{\theta}$	-0.30	-0.26	-0.38	-0.28
Var.	0.16	0.02	0.07	0.01	Var.	0.06	0.01	0.05	0.01
Bias ²	—	-4.20	0.01	-3.56	Bias ²	—	2.21	0.03	1.29
GFIC	0.16	-4.18	0.08	-3.54	GFIC	0.06	2.22	0.08	1.30
GFIC+	0.16	0.02	0.08	0.01	GFIC+	0.06	2.22	0.08	1.30

Table 3: Estimates and GFIC values for the price elasticity of demand for cigarettes example from Section 7 under four alternative specifications. Panel (a) presents results using data from 1975–1980 while Panel (b) presents results using data from 1975–1985. GFIC+ gives an alternative version of the GFIC in which a negative squared bias estimate is set equal to zero.

a key feature of the GFIC, namely that it implicitly considers the number of available time periods when selecting over parameter restrictions and moment conditions. Each column in Table 3 refers to a particular specification: LP, LS, P or S. The first row of each panel gives the associated estimate of the target parameter θ while the second gives the estimated asymptotic variance of $\sqrt{n}(\hat{\theta} - \theta_0)$, one of the two ingredients of the GFIC.¹⁴ Unsurprisingly the asymptotic variance decreases with the number of time periods available for estimation: for a given sample period LP and LS have a higher asymptotic variance than P and S, and all the estimates based on the 1975–1985 sample are more precise than their counterparts for the 1975–1980 sample. Moreover, for a given sample period, the estimators with a large instrument set show a lower asymptotic variance: LS is more precisely estimated than LP and S is more precisely estimated than P.

The third row of each panel gives our estimate of the squared asymptotic bias of the various estimators of θ . The “—” entry for the LP estimator indicates that the GFIC is constructed under the assumption that this specification has no asymptotic bias. Note that the squared bias estimator is *negative* for LS and S in the 1975–1980 sample. This occurs when the second term in the estimate of the bias matrix \hat{B} from Corollary 3.5, namely $\hat{\Psi}\hat{\Omega}\hat{\Psi}'$, is larger than the first. Accordingly, we consider two alternative ways of constructing the GFIC from the asymptotic variance and squared bias estimates. The first, labelled “GFIC,” simply adds squared bias and variance. The second, labelled “GFIC+,” first *truncates* a negative squared bias estimate to zero, and then adds the result to the variance estimate. For the 1975–1980 sample period we find no evidence of appreciable bias in any of the three “suspect” specifications: LS, P, and LP. In contrast, each of these has a substantially small asymptotic variance than LP, so we would select either LS or S depending on whether we

¹⁴We estimate the asymptotic variance matrix as in Baltagi et al. (2000).

prefer to use GFIC or GFIC+.¹⁵ In the 1975–1985 sample, the situation changes drastically. Over this longer time period, the relative advantage of LS, P and S over LP in asymptotic variance decreases substantially and we find evidence of substantial bias in both the LS and S specifications. It appears that over these additional time periods, the assumption that $\ln P_{it}$ is strictly exogenous fails. Interestingly the difference in GFIC values between LP and P in this longer sample is negligible. As P has a lower GFIC value than LP in the 1975–1980 sample, it appears that accounting for habit persistence is relatively unimportant in estimating the short-run price elasticity of cigarette demand.

8 Conclusion

This paper has introduced the GFIC, a proposal to choose moment conditions and parameter restrictions based on the quality of the estimates they provide. The GFIC performs well in simulations for our panel data examples, as does our proposed averaging estimator that combines the random and fixed effects specifications. While we focus here on applications to panel data, the GFIC can be applied to any GMM problem in which a minimal set of correctly specified moment conditions identifies an unrestricted model. A possible extension of this work would be to consider risk functions other than MSE, by analogy to [Claeskens et al. \(2006\)](#) and [Claeskens and Hjort \(2008\)](#). Another possibility would be to derive a version of the GFIC for GEL estimators. Although first-order equivalent to GMM, GEL estimators often exhibit superior finite-sample properties and may thus improve the quality of the selection criterion ([Newey and Smith, 2004](#)).

References

- Anderson, T., Hsiao, C., 1982. Formulation and estimation of dynamic models using panel data. *Journal of Econometrics* 18, 47–82.
- Andrews, D. W. K., Lu, B., 2001. Consistent model and moment selection procedures for GMM estimation with application to dynamic panel data models. *Journal of Econometrics* 101, 123–164.
- Arellano, M., Bond, S., 1991. Some tests of specification for panel data: Monte carlo evidence and an application to employment equations. *The Review of Economic Studies* 58 (2), 277–297.
- Baltagi, B. H., Griffin, J. M., Xiong, W., 2000. To pool or not to pool: Homogeneous versus heterogeneous estimators applied to cigarette demand. *The Review of Economics and Statistics* 82, 117–126.
- Caner, M., 2009. Lasso-type GMM estimator. *Econometric Theory* 25, 270–290.

¹⁵When GFIC and GFIC+ disagree, we prefer GFIC+ for the same reason that the positive-part Stein estimator is preferred to the “plain-vanilla” Stein estimator.

- Claeskens, G., Croux, C., Kerckhoven, J. V., 2006. Variable selection for logistic regression using a prediction-focused information criterion. *Biometrics* 62, 972–979.
- Claeskens, G., Hjort, N. L., 2003. The focused information criterion. *Journal of the American Statistical Association* 98 (464), 900–945.
- Claeskens, G., Hjort, N. L., 2008. Minimizing average risk in regression models. *Econometric Theory* 24, 493–527.
- DiTraglia, F. J., 2016. Using invalid instruments on purpose: Focused moment selection and averaging for GMM. *Journal of Econometrics* 195 (2), 187–208.
- Guggenberger, P., 2010. The impact of a Hausman pretest on the size of a hypothesis test: The panel data case. *Journal of Econometrics* 156, 337–343.
- Hall, A. R., Peixe, F. P., 2003. A consistent method for the selection of relevant instruments. *Econometric Reviews* 22, 269–287.
- Hansen, B. E., 2016. Efficient shrinkage in parametric models. *Journal of Econometrics* 190, 115–132.
- Hansen, B. E., 2017. Stein-like 2sls estimator. *Econometric Reviews*, 1–13.
- Hjort, N. L., Claeskens, G., 2003. Frequentist model average estimators. *Journal of the American Statistical Association* 98 (464), 879–899.
- Hong, H., Preston, B., Shum, M., 2003. Generalized empirical likelihood-based model selection criteria for moment condition models. *Econometric Theory* 19, 923–943.
- Judge, G. G., Mittelhammer, R. C., 2007. Estimation and inference in the case of competing sets of estimating equations. *Journal of Econometrics* 138, 513–531.
- Lai, T. L., Small, D. S., Liu, J., 2008. Statistical inference in dynamic panel data models. *Journal of Statistical Planning and Inference* 138, 2763–2776.
- Leeb, H., Pötscher, B. M., 2008. Sparse estimators and the oracle property, or the return of Hodges’ estimator. *Journal of Econometrics* 142, 201–211.
- Newey, W. K., Smith, R. J., 2004. Higher order properties of gmm and generalized empirical likelihood estimators. *Econometrica* 72 (1), 219–255.
- Pesaran, M. H., Shin, Y., Smith, R. P., 1999. Pooled mean group estimation of dynamic heterogeneous panels. *Journal of the American Statistical Association* 94, 621–634.
- Pesaran, M. H., Smith, R. P., 1995. Estimating long-run relationships from dynamic heterogeneous panels. *Journal of Econometrics* 68, 79–113.
- Roodman, D., 2009. A note on the theme of too many instruments. *Oxford Bulletin of Economics and Statistics* 71 (1), 135–158.
- Schorfheide, F., 2005. VAR forecasting under misspecification. *Journal of Econometrics* 128, 99–136.

Smith, R. J., 1992. Non-nested tests for competing models estimated by generalized method of moments. *Econometrica* 60 (4), 973–980.

Swamy, P. A. V. B., 1970. Efficient inference in a random coefficient regression model. *Econometrica* 38, 311–323.

Yang, Y., 2005. Can the strengths of AIC and BIC be shared? a conflict between model identification and regression estimation. *Biometrika* 92 (4), 937–950.

A Proofs

Proof of Theorem 2.1. By a mean-value expansion around (γ_0, θ_0) ,

$$\sqrt{n} \left(\hat{\beta}(b, c) - \beta_0^{(b)} \right) = -K(b, c) \Xi_c \sqrt{n} f_n(\gamma_0, \theta_0) + o_p(1)$$

and by assumption 2.2, $\sqrt{n} f_n(\gamma_0, \theta_0) - \sqrt{n} E[f(Z_{ni}, \gamma_0, \theta_0)] \xrightarrow{d} (\mathcal{N}'_g, \mathcal{N}'_h)'$. Now, by a mean-value expansion around γ_n ,

$$\begin{aligned} \sqrt{n} E[f(Z_{ni}, \gamma_0, \theta_0)] &= \sqrt{n} E[f(Z_{ni}, \gamma_n, \theta_0)] + \sqrt{n} \nabla_{\gamma'} E[f(Z_{ni}, \bar{\gamma}, \theta_0)] (\gamma_0 - \gamma_n) \\ &= \left(\begin{bmatrix} 0 \\ \tau \end{bmatrix} - \nabla_{\gamma'} E[f(Z_{ni}, \bar{\gamma}, \theta_0)] \delta \right) \rightarrow \left(\begin{bmatrix} 0 \\ \tau \end{bmatrix} - F_\gamma \delta \right). \end{aligned}$$

Hence, $\sqrt{n} f_n(\gamma_0, \theta_0) \xrightarrow{d} (\mathcal{N}'_g, \mathcal{N}'_h)' + (0', \tau')' - F_\gamma \delta$. The result follows by the continuous mapping theorem. \square

Proof of Corollary 2.1. Since Ξ_c picks out only the components corresponding to g , For the valid estimator,

$$\Xi_c \left(\begin{bmatrix} \mathcal{N}'_g \\ \mathcal{N}'_h \end{bmatrix} + \begin{bmatrix} 0 \\ \tau \end{bmatrix} - F_\gamma \delta \right) = \mathcal{N}'_g - G_\gamma \delta.$$

Thus, $\sqrt{n} (\hat{\beta}_v - \beta_0) \xrightarrow{d} -K_v (\mathcal{N}'_g - G_\gamma \delta)$. Finally, by the definition of a matrix inverse,

$$K_v G_\gamma \delta = \begin{bmatrix} G'_\gamma W_{gg} G_\gamma & G'_\gamma W_{gg} G_\theta \\ G'_\theta W_{gg} G_\gamma & G'_\theta W_{gg} G_\theta \end{bmatrix}^{-1} \begin{bmatrix} G'_\gamma W_{gg} G_\gamma \\ G'_\theta W_{gg} G_\gamma \end{bmatrix} \delta = \begin{bmatrix} 0 \\ \delta \end{bmatrix}.$$

\square

Proof of Corollary 3.2. For some $\bar{\gamma}$ between γ_0 and $\gamma_n = \gamma_0 + \delta/\sqrt{n}$.

$$\mu_n = \varphi(\gamma_0, \theta_0) + \nabla_\gamma \varphi(\bar{\gamma}, \theta_0)' (\gamma_n - \gamma_0) = \mu_0 + \nabla_\gamma \varphi(\bar{\gamma}, \theta_0)' \delta / \sqrt{n}$$

by a mean-value expansion. Hence, $\sqrt{n}(\mu_n - \mu_0) = \nabla_\gamma \varphi(\bar{\gamma}, \theta_0)' \delta \rightarrow \nabla_\gamma \varphi(\gamma_0, \theta_0)' \delta$. The result follows since $\sqrt{n}(\hat{\mu}(b, c) - \mu_n) = \sqrt{n}(\hat{\mu}(b, c) - \mu_0) - \sqrt{n}(\mu_n - \mu_0)$. \square

Proof of Corollary 3.3. The result follows from Corollaries 2.1 and 3.2 since,

$$\begin{aligned}
\sqrt{n}(\hat{\mu}_v - \mu_n) &\xrightarrow{d} \nabla_{\beta} \varphi'_0 \left\{ \begin{bmatrix} 0 \\ \delta \end{bmatrix} - K_v \mathcal{N}_g \right\} - \nabla_{\gamma} \varphi'_0 \delta \\
&= -\nabla_{\beta} \varphi'_0 K_v \mathcal{N}_g + \begin{bmatrix} \nabla_{\theta} \varphi'_0 & \nabla_{\gamma} \varphi'_0 \end{bmatrix} \begin{bmatrix} 0 \\ \delta \end{bmatrix} - \nabla_{\gamma} \varphi'_0 \delta \\
&= -\nabla_{\beta}(\gamma_0, \theta_0)' K_v \mathcal{N}_g.
\end{aligned}$$

□

Proof of Corollary 3.4. The result follows immediately from Corollary 2.1.

□

Proof of Lemma 3.1. By a mean-value expansion around $\beta_0 = (\gamma'_0, \theta'_0)'$,

$$\sqrt{n}h_n(\hat{\beta}_v) = \sqrt{n}h_n(\beta_0) + H\sqrt{n}(\hat{\beta}_v - \beta_0) + o_p(1).$$

Now, since

$$\sqrt{n}f_n(\gamma_0, \theta_0) \xrightarrow{d} \begin{bmatrix} \mathcal{N}_g \\ \mathcal{N}_h \end{bmatrix} + \begin{bmatrix} 0 \\ \tau \end{bmatrix} - \begin{bmatrix} G_{\gamma} \\ H_{\gamma} \end{bmatrix} \delta$$

we have $\sqrt{n}h_n(\gamma_0, \theta_0) \xrightarrow{d} \mathcal{N}_h + \tau - H_{\gamma}\delta$. Substituting,

$$\sqrt{n}h_n(\hat{\beta}_v) \xrightarrow{d} \mathcal{N}_h + \tau - H_{\gamma}\delta + H \left(-K_v \mathcal{N}_g + \begin{bmatrix} 0 \\ \delta \end{bmatrix} \right) = \tau - HK_v \mathcal{N}_g + \mathcal{N}_h.$$

□

Proof of Corollary 3.5. Define $(U', V')' = (\delta', \tau')' + \Psi(\mathcal{N}'_g, \mathcal{N}'_h)'$. By the Continuous Mapping Theorem and Theorem 3.1,

$$\begin{bmatrix} \hat{\delta} \\ \hat{\tau} \end{bmatrix} \begin{bmatrix} \hat{\delta}' & \hat{\tau}' \end{bmatrix} \xrightarrow{d} \begin{bmatrix} U \\ V \end{bmatrix} \begin{bmatrix} U' & V' \end{bmatrix}$$

The result follows since

$$\Psi \Omega \Psi' = Var \begin{bmatrix} U \\ V \end{bmatrix} = E \begin{bmatrix} UU' & UV' \\ VU' & VV' \end{bmatrix} - \begin{bmatrix} \delta\delta' & \delta\tau' \\ \tau\delta' & \tau\tau' \end{bmatrix}.$$

□

Proof of Theorem 4.1. To simplify the argument let $\xi = (\delta', \tau')'$ and $\mathcal{M} = \xi + \Psi \mathcal{N}$. Now, let $a(\xi^*)$ and $b(\xi^*)$ be quantiles of the distribution of Λ calculated under the assumption that the true bias parameter ξ is equal to ξ^* such that $P\{a(\xi^*) \leq \Lambda(\xi^*) \leq b(\xi^*)\} = 1 - \alpha_2$. Define

$$\begin{aligned}
a_{min}(\mathcal{M}) &= \min \{a(\xi^*) : \xi^* \in \mathcal{R}(\mathcal{M}|\alpha_1)\} \\
b_{max}(\mathcal{M}) &= \max \{b(\xi^*) : \xi^* \in \mathcal{R}(\mathcal{M}|\alpha_1)\} \\
\mathcal{R}(\mathcal{M}|\alpha_1) &= \left\{ \xi^* : \left[(\mathcal{M} - \xi^*)' (\Psi \Omega \Psi')^{-1} (\mathcal{M} - \xi^*) \right] \leq \chi^2_{p+q}(\alpha_1) \right\}
\end{aligned}$$

where $\chi^2_{p+q}(\alpha_1)$ is the $1 - \alpha_1$ quantile of a χ^2_{p+q} distribution. By Theorem 3.1 and Corollary 4.1,

$$P\{\mu_0 \in CI_{sim}\} \rightarrow P\{a_{min} \leq \Lambda(\delta, \tau) \leq b_{max}\}.$$

Let $A = \{\xi \in \mathcal{R}(\mathcal{M}|\alpha_1)\}$. By construction $P(A) = 1 - \alpha_1$. Decomposing $P\{a(\xi^*) \leq \Lambda(\xi^*) \leq b(\xi^*)\}$ into the sum of mutually exclusive events, we have

$$P[\{a(\xi^*) \leq \Lambda(\xi^*) \leq b(\xi^*)\} \cap A] + P[\{a(\xi^*) \leq \Lambda(\xi^*) \leq b(\xi^*)\} \cap A^c] = 1 - \alpha_2$$

for every $\xi \in \mathcal{R}(\mathcal{M}|\alpha_1)$. But since

$$P[\{a(\xi^*) \leq \Lambda(\xi^*) \leq b(\xi^*)\} \cap A^c] \leq P(A^c) = \alpha_1$$

we see that

$$P[\{a(\xi^*) \leq \Lambda(\xi^*) \leq b(\xi^*)\} \cap A] \geq 1 - \alpha_1 - \alpha_2$$

for every $\xi^* \in \mathcal{R}(\mathcal{M}|\alpha_1)$. By definition, if A occurs then $\xi \in \mathcal{R}(\mathcal{M}|\alpha_1)$ and hence

$$P[\{a(\xi) \leq \Lambda(\xi) \leq b(\xi)\} \cap A] \geq 1 - \alpha_1 - \alpha_2.$$

But when $\xi^* \in \mathcal{R}(\mathcal{M}|\alpha_1)$ we have $a_{min} \leq a(\xi)$ and $b(\xi) \leq b_{max}$. It follows that

$$\{a(\xi) \leq \Lambda(\xi) \leq b(\xi)\} \cap A \subseteq \{a_{min} \leq \Lambda(\xi) \leq b_{max}\}$$

and therefore

$$1 - \alpha_1 - \alpha_2 \leq P[\{a(\xi) \leq \Lambda(\xi) \leq b(\xi)\} \cap A] \leq P[a_{min} \leq \Lambda(\xi) \leq b_{max}].$$

□

Proof of Theorem 5.1. This proof is standard so we provide only a sketch. First, let $A_n = (n^{-1} \sum_{i=1}^n \mathbf{x}_i' \hat{\Omega}^{-1} \mathbf{x}_i)$, $B_n = (n^{-1} \sum_{i=1}^n \mathbf{x}_i' Q \mathbf{x}_i)$, and $C_n = T \hat{\sigma}_\alpha^2 + \hat{\sigma}_\varepsilon^2$. Now, expanding $\hat{\beta}_{FE}$, β_{RE} , and $\hat{\tau}$ and re-arranging

$$\begin{bmatrix} \sqrt{n}(\hat{\beta}_{RE} - \beta) \\ \sqrt{n}(\hat{\beta}_{FE} - \beta) \\ \hat{\tau} \end{bmatrix} = \begin{bmatrix} A_n^{-1} & 0 \\ 0 & B_n^{-1} \\ C_n & -C_n A_n B_n^{-1} \end{bmatrix} \begin{bmatrix} n^{-1/2} \sum_{i=1}^n \mathbf{x}_i' \hat{\Omega}^{-1} \mathbf{v}_i \\ n^{-1/2} \sum_{i=1}^n \mathbf{x}_i' Q \mathbf{v}_i \end{bmatrix}.$$

The result follows by applying a law of large numbers to A_n, B_n, C_n , and $\hat{\Omega}$ and the Lindeberg-Feller CLT jointly to $n^{-1/2} \sum_{i=1}^n \mathbf{x}_i' Q \mathbf{v}_i$ and $n^{-1/2} \sum_{i=1}^n \mathbf{x}_i' \hat{\Omega}^{-1} \mathbf{v}_i$. □

Proof of Theorem 5.2. This proof is standard, so we provide only a sketch. Expanding,

$$\begin{aligned} \sqrt{n} [\hat{\beta}(k, \cdot) - \beta] &= (0, \mathbf{0}'_{k-1}, \delta)' + \hat{Q}(k, \cdot) [n^{-1/2} Z'(k, \cdot) \Delta \mathbf{v}] \\ \sqrt{n} [\hat{\beta}(r, \cdot) - \beta_r] &= \hat{Q}(r, \cdot) \left\{ \delta [n^{-1} Z'(r, \cdot) L^k \Delta \mathbf{y}^+] + [n^{-1/2} Z'(r, \cdot) \Delta \mathbf{v}^+] \right\} \end{aligned}$$

The result follows, after some algebra, by applying the Lindeberg-Feller CLT to $n^{-1/2} Z'(k, \cdot) \Delta \mathbf{v}$ and $n^{-1/2} Z'(r, \cdot) \Delta \mathbf{v}^+$ and an appropriate law of large numbers to $n^{-1} Z'(r, \cdot) L^k \Delta \mathbf{y}^+$, where (\cdot) is either P or S depending on the instrument set used. □

Proof of Theorem 5.3. Expanding $\widehat{\beta}(k, P)$ from Equation 29

$$\begin{aligned} n^{-1/2} X' [\Delta \mathbf{y} - W(k) \widehat{\beta}(k, P)] &= n^{-1/2} X' [\Delta \mathbf{v} - W(k) \widehat{Q}(k, P) n^{-1} Z'(k, P) \Delta \mathbf{v}] \\ &= \begin{bmatrix} -n^{-1} X' W(k) \widehat{Q}(k, P) & I_{T-k-1} \end{bmatrix} \begin{bmatrix} n^{-1/2} Z'(k, P) \Delta \mathbf{v} \\ n^{-1/2} X' \Delta \mathbf{v} \end{bmatrix} \end{aligned}$$

using $\Delta \mathbf{y} = W(k) \beta_n + \Delta \mathbf{v}$. Similarly, expanding $\widehat{\beta}(k, P)$ from Equation 29,

$$\sqrt{n} [\widehat{\beta}(k, P)] = \begin{bmatrix} 0 & \mathbf{0}_{k-1} & \delta \end{bmatrix} + \widehat{Q}(k, P) n^{-1/2} Z'(k, P) \Delta \mathbf{v}$$

and since $\widehat{\delta}$ is \sqrt{n} times the k^{th} element of $\widehat{\beta}(k, P)$ and the k^{th} element β is zero, we have

$$\widehat{\delta} = \delta + \begin{bmatrix} 0 & \mathbf{0}'_{k-1} & 1 \end{bmatrix} \widehat{Q}(k, P) n^{-1/2} Z'(k, P) \Delta \mathbf{v}.$$

By a law of large numbers $\widehat{Q}(k, P) \rightarrow_p Q(k, P)$ and $n^{-1} X' W(k) \rightarrow_p \boldsymbol{\xi}' \otimes \boldsymbol{\iota}_{T-k-1}$, hence

$$\begin{bmatrix} \widehat{\delta} - \delta \\ \widehat{\tau} \end{bmatrix} = \Psi \begin{bmatrix} n^{-1/2} Z'(k, P) \Delta \mathbf{v} \\ n^{-1/2} X' \Delta \mathbf{v} \end{bmatrix} + o_p(1).$$

The result follows by applying the Lindeberg-Feller CLT jointly to $n^{-1/2} Z'(k, P) \Delta \mathbf{v}$ and $n^{-1/2} X' \Delta \mathbf{v}$ and noting that Π maps $[n^{-1/2} Z'(k, S) \Delta \mathbf{v}]$ to $\left[\{n^{-1/2} Z'(k, P) \Delta \mathbf{v}\}' \quad \{n^{-1/2} X' \Delta \mathbf{v}\}' \right]'$. \square

Proof of Theorem 5.4. Expanding the definitions of the OLS and mean-group estimators,

$$\begin{aligned} \sqrt{n}(\widehat{\beta}_{OLS} - \beta) &= \begin{bmatrix} (n^{-1} \sum_{i=1}^n \mathbf{x}'_i \mathbf{x}_i)^{-1} & (n^{-1} \sum_{i=1}^n \mathbf{x}'_i \mathbf{x}_i)^{-1} \end{bmatrix} \begin{bmatrix} n^{-1/2} \sum_{i=1}^n \mathbf{x}'_i \mathbf{x}_i \eta_i \\ n^{-1/2} \sum_{i=1}^n \mathbf{x}'_i \varepsilon_i \end{bmatrix} \\ \sqrt{n}(\widehat{\beta}_{MG} - \beta) &= n^{-1/2} \sum_{i=1}^n [\eta_i + (\mathbf{x}'_i \mathbf{x}_i)^{-1} \mathbf{x}'_i \varepsilon_i] \end{aligned}$$

and proceeding similarly for $\widehat{\tau}$,

$$\widehat{\tau} = \begin{bmatrix} 1 & 1 & -n^{-1} \sum_{i=1}^n \mathbf{x}'_i \mathbf{x}_i \end{bmatrix} \begin{bmatrix} n^{-1/2} \sum_{i=1}^n \mathbf{x}'_i \mathbf{x}_i \eta_i \\ n^{-1/2} \sum_{i=1}^n \mathbf{x}'_i \varepsilon_i \\ n^{-1/2} \sum_{i=1}^n \{\eta_i + (\mathbf{x}'_i \mathbf{x}_i)^{-1} \mathbf{x}'_i \varepsilon_i\} \end{bmatrix}.$$

The result follows, after some algebra, by a LLN and the Lindeberg-Feller CLT. \square

B Supplementary Simulation Results

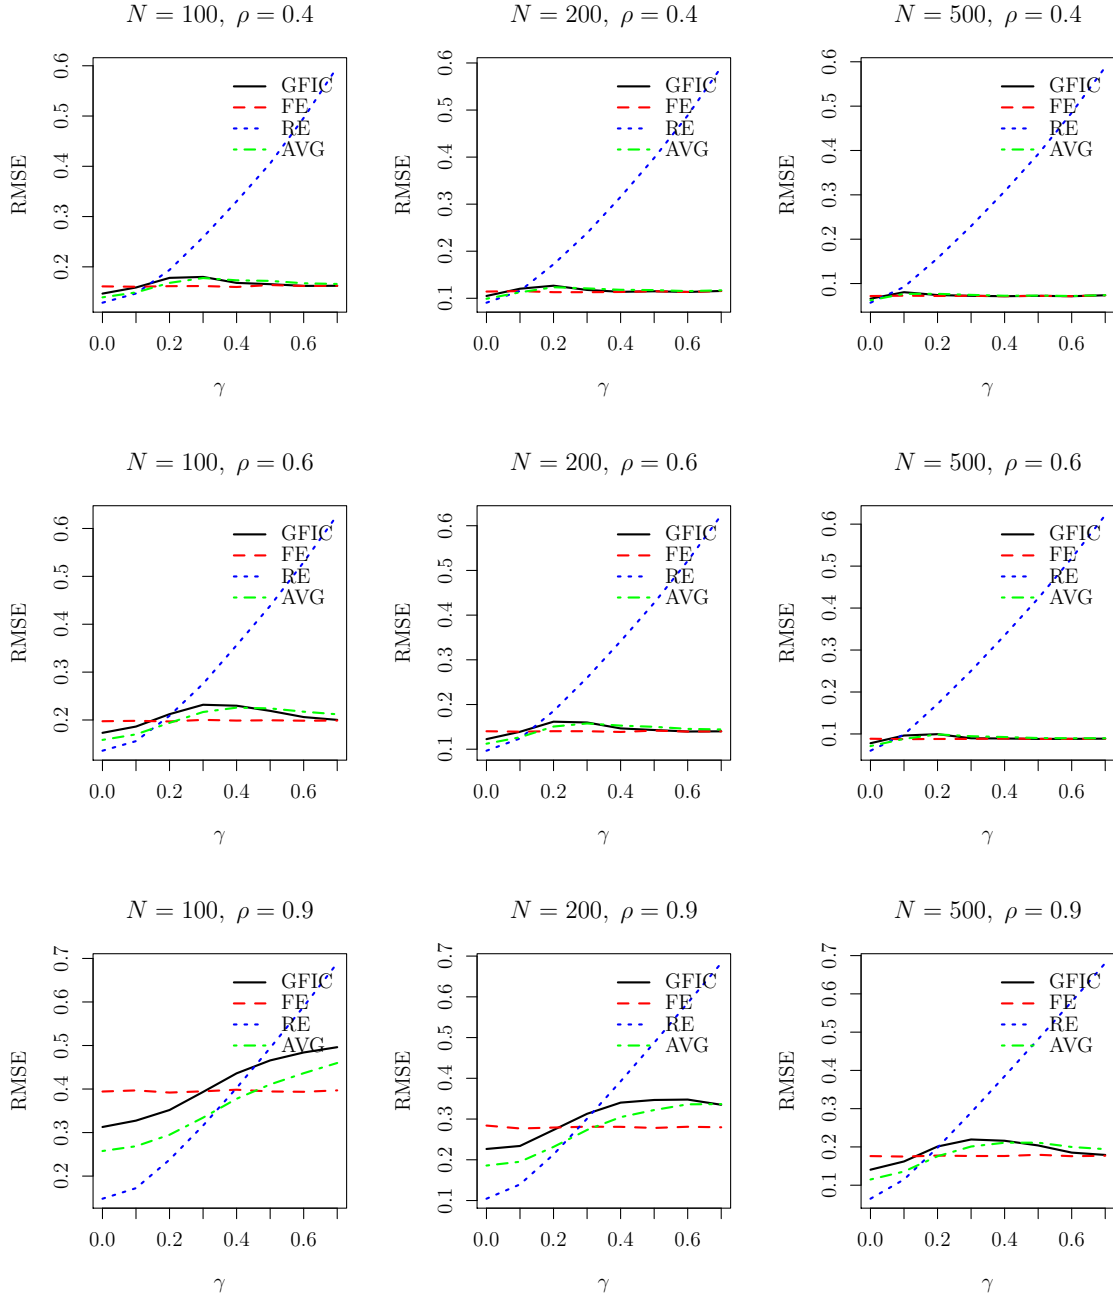


Figure 4: RMSE for Random vs. Fixed effects example from Section 6.1: $T = 5, \sigma_\varepsilon^2 = 2.5$




Figure 5: RMSE of GFIC relative to Oracle Estimator in the Simulation from Section 6.2.2




Figure 6: RMSE of GFIC relative to GMM-AIC in the Simulation from Section 6.2.2




Figure 7: RMSE of GFIC relative to GMM-BIC in the Simulation from Section 6.2.2




Figure 8: RMSE of GFIC relative to 5% Downward J-test in the Simulation from Section 6.2.2




Figure 9: RMSE of GFIC relative to 10% Downward J-test in the Simulation from Section 6.2.2




Figure 10: RMSE of GFIC relative to GMM-HQ in the Simulation from Section 6.2.2